

2.1 Process concept

Introduction

Overview

A *process* is a program in execution. A process will need certain resources—such as CPU time, memory, files, and I/O devices—to accomplish its task. These resources are typically allocated to the process while it is executing.

A process is the unit of work in most systems. Systems consist of a collection of processes: operating-system processes execute system code, and user processes execute user code. All these processes may execute concurrently.

Modern operating systems support processes having multiple *threads* of control. On systems with multiple hardware processing cores, these threads can run in parallel.

One of the most important aspects of an operating system is how it schedules threads onto available processing cores. Several choices for designing CPU schedulers are available to programmers.

Early computers allowed only one program to be executed at a time. This program had complete control of the system and had access to all the system's resources. In contrast, contemporary computer systems allow multiple programs to be loaded into memory and executed concurrently. This evolution required firmer control and more compartmentalization of the various programs; and these needs resulted in the notion of a **process**, which is a program in execution. A process is the unit of work in a modern computing system.

The more complex the operating system is, the more it is expected to do on behalf of its users. Although its main concern is the execution of user programs, it also needs to take care of various system tasks that are best done in user space, rather than within the kernel. A system therefore consists of a collection of processes, some executing user code, others executing operating system code. Potentially, all these processes can execute concurrently, with the CPU (or CPUs) multiplexed among them. In this chapter, you will read about what processes are, how they are represented in an operating system, and how they work.

Chapter objectives

-
- Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system.
 - Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations.
 - Describe and contrast interprocess communication using shared memory and message passing.
 - Design programs that use pipes and POSIX shared memory to perform interprocess communication.
 - Describe client-server communication using sockets and remote procedure calls.
 - Design kernel modules that interact with the Linux operating system.
-

Section glossary

process: A program loaded into memory and executing.

Process concept

A question that arises in discussing operating systems involves what to call all the CPU activities. Early computers were batch systems that executed **jobs**, followed by the emergence of time-shared systems that ran **user programs**, or **tasks**. Even on a single-user system, a user may be able to run several programs at one time: a word processor, a web browser, and an e-mail package. And even if a computer can execute only one program at a time, such as on an embedded device that does not support multitasking, the operating system may need to support its own internal programmed activities, such as memory management. In many respects, all these activities are similar, so we call all of them **processes**.

Although we personally prefer the more contemporary term **process**, the term **job** has historical significance, as much of operating system theory and terminology was developed during a time when the major activity of operating systems was job processing. Therefore, in some appropriate instances we use **job** when describing the role of the operating system. As an example, it would be misleading to avoid the use of commonly accepted terms that include the word **job** (such as **job scheduling**) simply because **process** has superseded **job**.


The process

Informally, as mentioned earlier, a process is a program in execution. The status of the current activity of a process is represented by the value of the **program counter** and the contents of the processor's registers.

The memory layout of a process is typically divided into multiple sections, and is shown in Figure [2.1.1](#). These sections include:

- **Text section**—the executable code
- **Data section**—global variables
- **Heap section**—memory that is dynamically allocated during program run time
- **Stack section**—temporary data storage when invoking functions (such as function parameters, return addresses, and local variables)

Figure 2.1.1: Layout of a process in memory.



Notice that the sizes of the text and data sections are fixed, as their sizes do not change during program run time. However, the stack and heap sections can shrink and grow dynamically during program execution. Each time a function is called, an **activation record** containing function parameters, local variables, and the return address is pushed onto the stack; when control is returned from the function, the activation record is popped from the stack. Similarly, the heap will grow as memory is dynamically allocated, and will shrink when memory is returned to the system. Although the stack and heap sections grow **toward** one another, the operating system must ensure they do not **overlap** one another.

We emphasize that a program by itself is not a process. A program is a **passive** entity, such as a file containing a list of instructions stored on disk (often called an **executable file**). In contrast, a process is an **active** entity, with a program counter specifying the next instruction to execute and a set of associated resources. A program becomes a process when an executable file is loaded into memory. Two common techniques for loading executable files are double-clicking an icon representing the executable file and entering the name of the executable file on the command line (as in **prog.exe** or **a.out**).

Although two processes may be associated with the same program, they are nevertheless considered two separate execution sequences. For instance, several users may be running different copies of the mail program, or the same user may invoke many copies of the web browser program. Each of these is a separate process; and although the text sections are equivalent, the data, heap, and stack sections vary. It is also common to have a process that spawns many processes as it runs. We discuss such matters in Section [2.4](#).

Note that a process can itself be an execution environment for other code. The Java programming environment provides a good example. In most circumstances, an executable Java program is executed within the Java virtual machine (JVM). The JVM executes as a process that interprets the loaded Java code and takes actions (via native machine instructions) on behalf of that code. For example, to run the compiled Java program `Program.class`, we would enter

```
java Program
```

The command `java` runs the JVM as an ordinary process, which in turns executes the Java program `Program` in the virtual machine. The concept is the same as simulation, except that the code, instead of being written for a different instruction set, is written in the Java language.

Process state


As a process executes, it changes **state**. The state of a process is defined in part by the current activity of that process. A process may be in one of the following states:

- **New.** The process is being created.
- **Running.** Instructions are being executed.
- **Waiting.** The process is waiting for some event to occur (such as an I/O completion or reception of a signal).
- **Ready.** The process is waiting to be assigned to a processor.
- **Terminated.** The process has finished execution.

Memory layout of a C program

The figure shown below illustrates the layout of a C program in memory, highlighting how the different sections of a process relate to an actual C program. This figure is similar to the general concept of a process in memory as shown in Figure [2.1.1](#), with a few differences:

- The global data section is divided into different sections for (a) initialized data and (b) uninitialized data.
- A separate section is provided for the `argc` and `argv` parameters passed to the `main( )` function.



The GNU **size** command can be used to determine the size (in bytes) of some of these sections. Assuming the name of the executable file of the above C program is **memory**, the following is the output generated by entering the command **size memory**:


text	data	bss	dec	hex	filename
1158	284	8	1450	5aa	memory


The **data** field refers to initialized data, and **bss** refers to uninitialized data. (**bss** is a historical term referring to **block started by symbol**.) The **dec** and **hex** values are the sum of the three sections represented in decimal and hexadecimal, respectively.

These names are arbitrary, and they vary across operating systems. The states that they represent are found on all systems, however. Certain operating systems also more finely delineate process states. It is important to realize that only one process can be *running* on any processor core at any instant. Many processes may be *ready* and *waiting*, however. The state diagram corresponding to these states is presented in the animation below.

PARTICIPATION ACTIVITY

2.1.1: Process States.





Animation content:

Step 1: A new process is created. The new process is moved to the ready state. Step 2: The process is scheduled onto a core that is available and running. If an I/O request or an event request occurs, the process is moved to the waiting state. Step 3: Once the I/O request or event request is completed, the process is moved back to the ready state. Step 4: If the process is running and the core is needed, for example by an interrupt, the process is moved back to the ready state. Step 5: The cycle continues until the process exits, fails, or is terminated. The process then moves to the terminated state.

Animation captions:

1. A new process is created. The new process is moved to the ready state.
2. The process is scheduled onto a core that is available and running. If an I/O request or an event request occurs, the process is moved to the waiting state.
3. Once the I/O request or event request is completed, the process is moved back to the ready state.
4. If the process is running and the core is needed, for example by an interrupt, the process is moved back to the ready state.
5. The cycle continues until the process exits, fails, or is terminated. The process then moves to the terminated state.

PARTICIPATION ACTIVITY

2.1.2: Process State Transitions.



Match each state transition to the corresponding event.

How to use this tool


Running -> Waiting

Running -> Ready

Ready -> Running

Waiting -> Ready

Running -> Terminated



	Process exits.
	Process is selected to run on a processing core.
	A timeout has occurred.
	Process needs to wait for an event (such as completion of an I/O operation).
	The event for which the process was waiting has occurred.
Reset	

Process control block

Each process is represented in the operating system by a **process control block** (*PCB*)—also called a **task control block**. A PCB is shown in Figure [2.1.2](#). It contains many pieces of information associated with a specific process, including these:

- **Process state.** The state may be new, ready, running, waiting, halted, and so on.
- **Program counter.** The counter indicates the address of the next instruction to be executed for this process.
- **CPU registers.** The registers vary in number and type, depending on the computer architecture. They include accumulators, index registers, stack pointers, and general-purpose registers, plus any condition-code information. Along with the program counter, this state information must be saved when an interrupt occurs, to allow the process to be continued correctly afterward when it is rescheduled to run.
- **CPU-scheduling information.** This information includes a process priority, pointers to scheduling queues, and any other scheduling parameters. (Chapter CPU Scheduling describes process scheduling.)
- **Memory-management information.** This information may include such items as the value of the base and limit registers and the page tables, or the segment tables, depending on the memory system used by the operating system (Chapter Main Memory).
- **Accounting information.** This information includes the amount of CPU and real time used, time limits, account numbers, job or process numbers, and so on.
- **I/O status information.** This information includes the list of I/O devices allocated to the process, a list of open files, and so on.

In brief, the PCB simply serves as the repository for all the data needed to start, or restart, a process, along with some accounting data.

Figure 2.1.2: Process control block (PCB)



Threads

The process model discussed so far has implied that a process is a program that performs a single **thread** of execution. For example, when a process is running a word-processor program, a single thread of instructions is being executed. This single thread of control allows the process to perform only one task at a time. Thus, the user cannot simultaneously type in characters and run the spell checker. Most modern operating systems have extended the process concept to allow a process to have multiple threads of execution and thus to perform more than one task at a time. This feature is especially beneficial on multicore systems, where multiple threads can run in parallel. A multithreaded word processor could, for example, assign one thread to manage user input while another thread runs the spell checker. On systems that support threads, the PCB is expanded to include information for each thread. Other changes throughout the system are also needed to support threads. Chapter Threads & Concurrency explores threads in detail.

PARTICIPATION
ACTIVITY

2.1.3: Layout of a process in memory.

Organize the layout of a process in memory.

How to use this tool ▼

Unused blocks

↓

Data Section

↑

Text Section

Heap Section

Stack Section

Proof

Reset

Move blocks here

0 of 6 blocks correct

**PARTICIPATION
ACTIVITY**

2.1.4: Section review questions.



1) The convention in modern operating systems and compilers is that the ____ of a process contains temporary data such as function parameters, return addresses, and local variables.



- ☐ text section
- ☐ data section
- ☐ stack

2) A process control block ____.



- ☐ contains a process' states
- ☐ stores the executable code the process will run
- ☐ determines which process is to be run next

3) A process may transition to the Ready state by which of the following actions?



- ☐ While the process is waiting for completion of an I/O event
- ☐ An interrupt has occurred or an I/O operation the process is waiting for has completed.
- ☐ The process is dispatched by the scheduler to run on a CPU core.

Section glossary

job: A set of commands or processes executed by a batch system.

user programs: User-level programs, as opposed to system programs.

task: A process, a thread activity, or, generally, a unit of computation on a computer.

process: A program loaded into memory and executing.

program counter: A CPU register indicating the main memory location of the next instruction to load and execute.

text section: The executable code of a program or process.

data section: The data part of a program or process; it contains global variables.

heap section: The section of process memory that is dynamically allocated during process run time; it stores temporary variables.

stack section: The section of process memory that contains the stack; it contains activation records and other temporary data.

activation record: A record created when a function or subroutine is called; added to the stack by the call and removed when the call returns. Contains function parameters, local variables, and the return address.

executable file: A file containing a program that is ready to be loaded into memory and executed.

state: The condition of a process, including its current activity as well as its associated memory and disk contents.

process control block: A per-process kernel data structure containing many pieces of information associated with the process.

task control block: A per-process kernel data structure containing many pieces of information associated with the process.

thread: A process control structure that is an execution location. A process with a single thread executes only one task at a time, while a multithreaded process can execute a task per thread.

2.2 Process scheduling


The objective of multiprogramming is to have some process running at all times so as to maximize CPU utilization. The objective of time sharing is to switch a CPU core among processes so frequently that users can interact with each program while it is running. To meet these objectives, the **process scheduler** selects an available process (possibly from a set of several available processes) for program execution on a core. Each CPU core can run one process at a time. For a system with a single CPU core, there will never be more than one process running at a time, whereas a multicore system can run multiple processes at one time. If there are more processes than cores, excess processes will have to wait until a core is free and can be rescheduled. The number of processes currently in memory is known as the **degree of multiprogramming**.

Process representation in linux

The process control block in the Linux operating system is represented by the C structure **task_struct**, which is found in the `<include/linux/sched.h>` include file in the kernel source-code directory. This structure contains all the necessary information for representing a process, including the state of the process, scheduling and memory-management information, list of open files, and pointers to the process's parent and a list of its children and siblings. (A process's **parent** is the process that created it; its **children** are any processes that it creates. Its **siblings** are children with the same parent process.) Some of these fields include:

```
long state;                /* state of the process */
struct sched_entity se;     /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children;  /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;       /* address space */
```

For example, the state of a process is represented by the field **long state** in this structure. Within the Linux kernel, all active processes are represented using a doubly linked list of **task_struct**. The kernel maintains a pointer—**current**—to the process currently executing on the system, as shown below:



As an illustration of how the kernel might manipulate one of the fields in the `task_struct` for a specified process, let's assume the system would like to change the state of the process currently running to the value `new_state`. If `current` is a pointer to the process currently executing, its state is changed with the following:

```
current->state = new_state;
```


Balancing the objectives of multiprogramming and time sharing also requires taking the general behavior of a process into account. In general, most processes can be described as either I/O bound or CPU bound. An **I/O-bound process** is one that spends more of its time doing I/O than it spends doing computations. A **CPU-bound process**, in contrast, generates I/O requests infrequently, using more of its time doing computations.

Scheduling queues

As processes enter the system, they are put into a **ready queue**, where they are ready and waiting to execute on a CPU's core. This queue is generally stored as a linked list; a ready-queue header contains pointers to the first PCB in the list, and each PCB includes a pointer field that points to the next PCB in the ready queue.


The system also includes other queues. When a process is allocated a CPU core, it executes for a while and eventually terminates, is interrupted, or waits for the occurrence of a particular event, such as the completion of an I/O request. Suppose the process makes an I/O request to a device such as a disk. Since devices run significantly slower than processors, the process will have to wait for the I/O to become available. Processes that are waiting for a certain event to occur — such as completion of I/O — are placed in a **wait queue** (Figure [2.2.1](#))

Figure 2.2.1: The ready queue and wait queues.



A common representation of process scheduling is a **queueing diagram**, such as that in Figure 2.2.2. Two types of queues are present: the ready queue and a set of wait queues. The circles represent the resources that serve the queues, and the arrows indicate the flow of processes in the system.

Figure 2.2.2: Queueing-diagram representation of process scheduling.



A new process is initially put in the ready queue. It waits there until it is selected for execution, or **dispatched**. Once the process is allocated a CPU core and is executing, one of several events could occur:

- The process could issue an I/O request and then be placed in an I/O wait queue.
- The process could create a new child process and then be placed in a wait queue while it awaits the child's termination.
- The process could be removed forcibly from the core, as a result of an interrupt or having its time slice expire, and be put back in the ready queue.

In the first two cases, the process eventually switches from the waiting state to the ready state and is then put back in the ready queue. A process continues this cycle until it terminates, at which time it is removed from all queues and has its PCB and resources deallocated.

CPU scheduling

A process migrates among the ready queue and various wait queues throughout its lifetime. The role of the **CPU scheduler** is to select from among the processes that are in the ready queue and allocate a CPU core to one of them. The CPU scheduler must select a new process for the CPU frequently. An I/O-bound process may execute for only a few milliseconds before waiting for an I/O request. Although a CPU-bound process will require a CPU core for longer durations, the scheduler is unlikely to grant the core to a process for an extended period. Instead, it is likely designed to forcibly remove the CPU from a process and schedule another process to run. Therefore, the CPU scheduler executes at least once every 100 milliseconds, although typically much more frequently.

Some operating systems have an intermediate form of scheduling, known as **swapping**, whose key idea is that sometimes it can be advantageous to remove a process from memory (and from active contention for the CPU) and thus reduce the degree of multiprogramming. Later, the process can be reintroduced into memory, and its execution can be continued where it left off. This scheme is known as **swapping** because a process can be "swapped out" from memory to disk, where its current status is saved, and later "swapped in" from disk back to memory, where its status is restored. Swapping is typically only necessary when memory has been overcommitted and must be freed up. Swapping is discussed in the chapter Main Memory.

Context switch


As mentioned in Section Interrupts, interrupts cause the operating system to change a CPU core from its current task and to run a kernel routine. Such operations happen frequently on general-purpose systems. When an interrupt occurs, the system needs to save the current **context** of the process running on the CPU core so that it can restore that context when its processing is done, essentially suspending the process and then resuming it. The context is represented in the PCB of the process. It includes the value of the CPU registers, the process state (see the animation Process State Changes), and memory-management information. Generically, we perform a **state save** of the current state of the CPU core, be it in kernel or user mode, and then a **state restore** to resume operations.

Switching the CPU core to another process requires performing a state save of the current process and a state restore of a different process. This task is known as a **context switch** and is illustrated in the animation below. When a context switch occurs, the kernel saves the context of the old process in its PCB and loads the saved context of the new process scheduled to run. Context-switch time is pure overhead,

because the system does no useful work while switching. Switching speed varies from machine to machine, depending on the memory speed, the number of registers that must be copied, and the existence of special instructions (such as a single instruction to load or store all registers). A typical speed is a several microseconds.

PARTICIPATION ACTIVITY

2.2.1: Context switch from process to process.



Animation content:

Step 1: The system has two processes, P_0 and P_1 . P_1 is idle, P_0 is executing and executes a system call, or the system receives an interrupt. The operating system saves the state of P_0 in its PCB. Step 2: P_0 is now off core, so idle. The operating system restores the state of P_1 from its PCB and puts P_1 on the core. It continues execution. Step 3: P_1 execution is interrupted, the operating system saves its state to its PCB and restores the next process's state (P_0 in this case) to prepare it to continue execution. Step 4: The operating system puts P_0 onto the core, continuing its execution while P_1 is now idle, in ready state, waiting its turn.

Animation captions:

1. The system has two processes, P_0 and P_1 . P_1 is idle, P_0 is executing and executes a system call, or the system receives an interrupt. The operating system saves the state of P_0 in its PCB.

2. P_0 is now off core, so idle. The operating system restores the state of P_1 from its PCB and puts P_1 on the core. It continues execution.
3. P_1 execution is interrupted, the operating system saves its state to its PCB and restores the next process's state (P_0 in this case) to prepare it to continue execution.
4. The operating system puts P_0 onto the core, continuing its execution while P_1 is now idle, in ready state, waiting its turn.

Multitasking in mobile systems

Because of the constraints imposed on mobile devices, early versions of iOS did not provide user-application multitasking; only one application ran in the foreground while all other user applications were suspended. Operating-system tasks were multitasked because they were written by Apple and well behaved. However, beginning with iOS 4, Apple provided a limited form of multitasking for user applications, thus allowing a single foreground application to run concurrently with multiple background applications. (On a mobile device, the **foreground** application is the application currently open and appearing on the display. The **background** application remains in memory, but does not occupy the display screen.) The iOS 4 programming API provided support for multitasking, thus allowing a process to run in the background without being suspended. However, it was limited and only available for a few application types. As hardware for mobile devices began to offer larger memory capacities, multiple processing cores, and greater battery life, subsequent versions of iOS began to support richer functionality for multitasking with fewer restrictions. For example, the larger screen on iPad tablets allowed running two foreground apps at the same time, a technique known as **split-screen**.

Since its origins, Android has supported multitasking and does not place constraints on the types of applications that can run in the background. If an application requires processing while in the background, the application must use a **service**, a separate application component that runs on behalf of the background process. Consider a streaming audio application: if the application moves to the background, the service continues to send audio data to the audio device driver on behalf of the background application. In fact, the service will continue to run even if the background application is suspended. Services do not have a user interface and have a small memory footprint, thus providing an efficient technique for multitasking in a mobile environment.

Context-switch times are highly dependent on hardware support. For instance, some processors provide multiple sets of registers. A context switch here simply requires changing the pointer to the current register set. Of course, if there are more active processes than there are register sets, the system resorts to copying register data to and from memory, as before. Also, the more complex the operating system, the greater the amount of work that must be done during a context switch. As we will see in the chapter Main Memory, advanced memory-management techniques may require that extra data be switched with each context. For instance, the address space of the current process must be preserved as the space of the

next task is prepared for use. How the address space is preserved, and what amount of work is needed to preserve it, depend on the memory-management method of the operating system.

**PARTICIPATION
ACTIVITY**

2.2.2: Section review questions.



1) Which of the following is not an example of a wait queue?



- ☐ Ready queue
- ☐ Child termination queue
- ☐ I/O queue

2) _____ saves the state of the currently running process and restores the state of the next process to run.



- ☐ The CPU scheduler
- ☐ Swapping
- ☐ A context switch

Section glossary

process scheduler: A scheduler that selects an available process (possibly from a set of several processes) for execution on a CPU.

degree of multiprogramming: The number of processes in memory.

parent: In a tree data structure, a node that has one or more nodes connected below it.

children: In a tree data structure, nodes connected below another node.

siblings: In a tree data structure, child nodes of the same parent.

I/O-bound process: A process that spends more of its time doing I/O than doing computations

CPU-bound process: A process that spends more time executing on CPU than it does performing I/O.

ready queue: The set of processes ready and waiting to execute.

wait queue: In process scheduling, a queue holding processes waiting for an event to occur before they need to be put on CPU.

dispatched: Selected by the process scheduler to be executed next.

CPU scheduler: Kernel routine that selects a thread from the threads that are ready to execute and allocates a core to that thread.

swapping: Moving a process between main memory and a backing store. A process may be swapped out to free main memory temporarily and then swapped back in to continue execution.

context: When describing a process, the state of its execution, including the contents of registers, its program counter, and its memory context, including its stack and heap.

state save: Copying a process's context to save its state in order to pause its execution in preparation for putting another process on the CPU.

state restore: Copying a process's context from its saved location to the CPU registers in preparation for continuing the process's execution.

context switch: The switching of the CPU from one process or thread to another; requires performing a state save of the current process or thread and a state restore of the other.

foreground: Describes a process or thread that is interactive (has input directed to it), such as a window currently selected as active or a terminal window currently selected to receive input.

background: Describes a process or thread that is not currently interactive (has no interactive input directed to it), such as one not currently being used by a user. In the Grand Central Dispatch Apple OS scheduler, the scheduling class representing tasks that are not time sensitive and are not visible to the user.

split-screen: Running multiple foreground processes (e.g., on an iPad) but splitting the screen among the processes.

service: A software entity running on one or more machines and providing a particular type of function to calling clients. In Android, an application component with no user interface; it runs in the background while executing long-running operations or performing work for remote processes.

2.3 Operations on processes

The processes in most systems can execute concurrently, and they may be created and deleted dynamically. Thus, these systems must provide a mechanism for process creation and termination. In this section, we explore the mechanisms involved in creating processes and illustrate process creation on UNIX and Windows systems.


Process creation

During the course of execution, a process may create several new processes. As mentioned earlier, the creating process is called a parent process, and the new processes are called the children of that process. Each of these new processes may in turn create other processes, forming a **tree** of processes.

Most operating systems (including UNIX, Linux, and Windows) identify processes according to a unique **process identifier** (or **pid**), which is typically an integer number. The pid provides a unique value for each process in the system, and it can be used as an index to access various attributes of a process within the kernel.

Figure 2.3.1 illustrates a typical process tree for the Linux operating system, showing the name of each process and its pid. (We use the term **process** rather loosely in this situation, as Linux prefers the term **task** instead.) The **systemd** process (which always has a pid of 1) serves as the root parent process for all user processes, and is the first user process created when the system boots. Once the system has booted, the **systemd** process creates processes which provide additional services such as a web or print server, an **ssh** server, and the like. In Figure 2.3.1, we see two children of **systemd**—**logind** and **sshd**. The **logind** process is responsible for managing clients that directly log onto the system. In this example, a client has logged on and is using the **bash** shell, which has been assigned pid 8416. Using the **bash** command-line interface, this user has created the process **ps** as well as the **vim** editor. The **sshd** process is responsible for managing clients that connect to the system by using **ssh** (which is short for **secure shell**).

Figure 2.3.1: A tree of processes on a typical Linux system.



The init and systemd processes

Traditional UNIX systems identify the process **init** as the root of all child processes. **init** (also known as **System V init**) is assigned a pid of 1, and is the first process created when the system is booted. On a process tree similar to what is shown in Figure 2.3.1, **init** is at the root.

Linux systems initially adopted the System V **init** approach, but recent distributions have replaced it with **systemd**. As described in Section Process creation, **systemd** serves as the system's initial process, much the same as System V **init**; however it is much more flexible, and can provide more services, than **init**.

On UNIX and Linux systems, we can obtain a listing of processes by using the **ps** command. For example, the command

```
ps -el
```

will list complete information for all processes currently active in the system. A process tree similar to the one shown in Figure [2.3.1](#) can be constructed by recursively tracing parent processes all the way to the **systemd** process. (In addition, Linux systems provide the **pstree** command, which displays a tree of all processes in the system.)

In general, when a process creates a child process, that child process will need certain resources (CPU time, memory, files, I/O devices) to accomplish its task. A child process may be able to obtain its resources directly from the operating system, or it may be constrained to a subset of the resources of the parent process. The parent may have to partition its resources among its children, or it may be able to share some resources (such as memory or files) among several of its children. Restricting a child process to a subset of the parent's resources prevents any process from overloading the system by creating too many child processes.

In addition to supplying various physical and logical resources, the parent process may pass along initialization data (input) to the child process. For example, consider a process whose function is to display the contents of a file—say, **hw1.c**—on the screen of a terminal. When the process is created, it will get, as an input from its parent process, the name of the file **hw1.c**. Using that file name, it will open the file and write the contents out. It may also get the name of the output device. Alternatively, some operating systems pass resources to child processes. On such a system, the new process may get two open files, **hw1.c** and the terminal device, and may simply transfer the datum between the two.

When a process creates a new process, two possibilities for execution exist:

1. The parent continues to execute concurrently with its children.
2. The parent waits until some or all of its children have terminated.

There are also two address-space possibilities for the new process:

1. The child process is a duplicate of the parent process (it has the same program and data as the parent).
2. The child process has a new program loaded into it.

To illustrate these differences, let's first consider the UNIX operating system. In UNIX, as we've seen, each process is identified by its process identifier, which is a unique integer. A new process is created by the **fork()** system call. The new process consists of a copy of the address space of the original process. This mechanism allows the parent process to communicate easily with its child process. Both processes (the parent and the child) continue execution at the instruction after the **fork()**, with one difference: the return code for the **fork()** is zero for the new (child) process, whereas the (nonzero) process identifier of the child is returned to the parent.

After a `fork()` system call, one of the two processes typically uses the `exec()` system call to replace the process's memory space with a new program. The `exec()` system call loads a binary file into memory (replacing the memory image of the program containing the `exec()` system call) and starts its execution. In this manner, the two processes are able to communicate and then go their separate ways. The parent can then create more children; or, if it has nothing else to do while the child runs, it can issue a `wait()` system call to move itself off the ready queue until the termination of the child. Because the call to `exec()` overlays the process's address space with a new program, `exec()` does not return control unless an error occurs.

The C program shown in Figure 2.3.2 illustrates the UNIX system calls previously described. We now have two different processes running copies of the same program. The only difference is that the value of the variable `pid` for the child process is zero, while that for the parent is an integer value greater than zero (in fact, it is the actual pid of the child process). The child process inherits privileges and scheduling attributes from the parent, as well certain resources, such as open files. The child process then overlays its address space with the UNIX command `/bin/ls` (used to get a directory listing) using the `execlp()` system call (`execlp()` is one of many different versions of the `exec()` system call). The parent waits for the child process to complete with the `wait()` system call. When the child process completes (by either implicitly or explicitly invoking `exit()`), the parent process resumes from the call to `wait()`, where it completes using the `exit()` system call. This is also illustrated in Figure 2.3.3.

Figure 2.3.2: Creating a separate process using the UNIX `fork()` system call.

```
#include <sys/types.h>
#include <sys/wait.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;


    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```


Copyrights 02/10/26 13:24 2150642
Adrian Tull
CS-510-10435.202617-1

Figure 2.3.3: Process creation using the fork() system call.



PARTICIPATION ACTIVITY

2.3.1: Process creation using the fork() system call.



```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0 { /* error occurred */
        fprintf(stderr, "Fork failed");
        return 1;
    }
    else if (pid == 0) { /*child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
  
```

Animation content:

Static figure: The static figure shows a fragment of source code, as well as a diagram showing how a process is created using the "fork()" system call. The diagram begins with the parent process running the statement "pid = fork()" in which it creates a child process using the fork() system call and assigns "pid" the value returned by the system call. In the child process the value of "pid" is 0, so it runs "exec()". The child process then terminates by calling "exit()" system call. In the parent process the value of "pid" is greater than 0, so it simply runs "wait()" until the child process completes. Lastly, once the child process terminates, the parent process returns from "wait()" and resumes. The source code shown alongside the diagram is shown below:

Begin C code:

```
include <sys/types.h>
include <stdio.h>
include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) {
        fprintf(stderr, "Fork failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

End C code.

Animation captions:

1. Parent calls the fork() system call which creates the child process. The child is a copy of the parent process which means there are two processes that return from the fork() system call – the parent and child processes.
2. The fork() system call returns 0 to identify the child process. In this case, the child process will run the exec() system call which runs the "ls" command.
3. The fork() system call returns a value > 0 to the parent and the parent calls the wait() system call waiting for the child process to terminate.
4. The child process will terminate by calling the exit() system call.
5. Once the child process terminates, the parent process returns from wait() and resumes.

Of course, there is nothing to prevent the child from **not** invoking **exec ()** and instead continuing to execute as a copy of the parent process. In this scenario, the parent and child are concurrent processes

running the same code instructions. Because the child is a copy of the parent, each process has its own copy of any data.

As an alternative example, we next consider process creation in Windows. Processes are created in the Windows API using the `CreateProcess()` function, which is similar to `fork()` in that a parent creates a new child process. However, whereas `fork()` has the child process inheriting the address space of its parent, `CreateProcess()` requires loading a specified program into the address space of the child process at process creation. Furthermore, whereas `fork()` is passed no parameters, `CreateProcess()` expects no fewer than ten parameters.

The C program shown in Figure 2.3.4 illustrates the `CreateProcess()` function, which creates a child process that loads the application `mspaint.exe`. We opt for many of the default values of the ten parameters passed to `CreateProcess()`. Readers interested in pursuing the details of process creation and management in the Windows API are encouraged to consult the bibliographical notes at the end of this chapter.

**PARTICIPATION
ACTIVITY**

2.3.2: Creating a child process using the `fork()` system call.

 **Full screen**



Complete the program that creates a child process using the `fork()` system call where the output is ordered by the child process first printing "I am the child" followed by the parent printing "I am the parent". Not all code sections may need to be used.

How to use this tool 

Unused

```
if (pid == 0)
{
if (pid > 0)
{
printf("I am the child.\n");
}
}

pid_t pid;
pid = fork();

else
{
printf("I am the parent.\n");
wait(NULL);
```

main.cpp

[Load default template...](#)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
int main()
{
    return 0;
}
```

©zyBooks 02/10/26 18:24 2150642
Adrian Tull
CS-510-10435-202617-1

Check

Figure 2.3.4: Creating a separate process using the Windows API.

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
}
```

©zyBooks 02/10/26 18:24 2150642
Adrian Tull
CS-510-10435-202617-1

```

/* parent will wait for the child to complete */
WaitForSingleObject(pi.hProcess, INFINITE);
printf("Child Complete");

/* close handles */
CloseHandle(pi.hProcess);
CloseHandle(pi.hThread);
}

```

The two parameters passed to the **CreateProcess()** function are instances of the **STARTUPINFO** and **PROCESS_INFORMATION** structures. **STARTUPINFO** specifies many properties of the new process, such as window size and appearance and handles to standard input and output files. The **PROCESS_INFORMATION** structure contains a handle and the identifiers to the newly created process and its thread. We invoke the **ZeroMemory()** function to allocate memory for each of these structures before proceeding with **CreateProcess()**.

The first two parameters passed to **CreateProcess()** are the application name and command-line parameters. If the application name is **NULL** (as it is in this case), the command-line parameter specifies the application to load. In this instance, we are loading the Microsoft Windows **mspaint.exe** application. Beyond these two initial parameters, we use the default parameters for inheriting process and thread handles as well as specifying that there will be no creation flags. We also use the parent's existing environment block and starting directory. Last, we provide two pointers to the **STARTUPINFO** and **PROCESS_INFORMATION** structures created at the beginning of the program. In Figure [2.3.2](#), the parent process waits for the child to complete by invoking the **wait()** system call. The equivalent of this in Windows is **WaitForSingleObject()**, which is passed a handle of the child process—**pi.hProcess**—and waits for this process to complete. Once the child process exits, control returns from the **WaitForSingleObject()** function in the parent process.

Process termination

A process terminates when it finishes executing its final statement and asks the operating system to delete it by using the **exit()** system call. At that point, the process may return a status value (typically an integer) to its waiting parent process (via the **wait()** system call). All the resources of the process—including physical and virtual memory, open files, and I/O buffers—are deallocated and reclaimed by the operating system.

Termination can occur in other circumstances as well. A process can cause the termination of another process via an appropriate system call (for example, **TerminateProcess()** in Windows). Usually, such a system call can be invoked only by the parent of the process that is to be terminated. Otherwise, a user—or a misbehaving application—could arbitrarily kill another user's processes. Note that a parent needs to know the identities of its children if it is to terminate them. Thus, when one process creates a new process, the identity of the newly created process is passed to the parent.

A parent may terminate the execution of one of its children for a variety of reasons, such as these:

- The child has exceeded its usage of some of the resources that it has been allocated. (To determine whether this has occurred, the parent must have a mechanism to inspect the state of its children.)

- The task assigned to the child is no longer required.
- The parent is exiting, and the operating system does not allow a child to continue if its parent terminates.

Some systems do not allow a child to exist if its parent has terminated. In such systems, if a process terminates (either normally or abnormally), then all its children must also be terminated. This phenomenon, referred to as **cascading termination**, is normally initiated by the operating system.

To illustrate process execution and termination, consider that, in Linux and UNIX systems, we can terminate a process by using the `exit()` system call, providing an exit status as a parameter:

```
/* exit with status 1 */
exit(1);
```

In fact, under normal termination, `exit()` will be called either directly (as shown above) or indirectly, as the C run-time library (which is added to UNIX executable files) will include a call to `exit()` by default.

A parent process may wait for the termination of a child process by using the `wait()` system call. The `wait()` system call is passed a parameter that allows the parent to obtain the exit status of the child. This system call also returns the process identifier of the terminated child so that the parent can tell which of its children has terminated:

```
pid_t pid;
int status;

pid = wait(&status);
```

When a process terminates, its resources are deallocated by the operating system. However, its entry in the process table must remain there until the parent calls `wait()`, because the process table contains the process's exit status. A process that has terminated, but whose parent has not yet called `wait()`, is known as a **zombie** process. All processes transition to this state when they terminate, but generally they exist as zombies only briefly. Once the parent calls `wait()`, the process identifier of the zombie process and its entry in the process table are released.

Now consider what would happen if a parent did not invoke `wait()` and instead terminated, thereby leaving its child processes as **orphans**. Traditional UNIX systems addressed this scenario by assigning the `init` process as the new parent to orphan processes. (Recall from Section Process creation that `init` serves as the root of the process hierarchy in UNIX systems.) The `init` process periodically invokes `wait()`, thereby allowing the exit status of any orphaned process to be collected and releasing the orphan's process identifier and process-table entry.

Although most Linux systems have replaced `init` with `systemd`, the latter process can still serve the same role, although Linux also allows processes other than `systemd` to inherit orphan processes and manage their termination.

Android process hierarchy

Because of resource constraints such as limited memory, mobile operating systems may have to terminate existing processes to reclaim limited system resources. Rather than terminating an arbitrary process, Android has identified an **importance hierarchy** of processes, and when the system must terminate a process to make resources available for a new, or more important, process, it terminates processes in order of increasing importance. From most to least important, the hierarchy of process classifications is as follows:

- **Foreground process**—The current process visible on the screen, representing the application the user is currently interacting with
- **Visible process**—A process that is not directly visible on the foreground but that is performing an activity that the foreground process is referring to (that is, a process performing an activity whose status is displayed on the foreground process)
- **Service process**—A process that is similar to a background process but is performing an activity that is apparent to the user (such as streaming music)
- **Background process**—A process that may be performing an activity but is not apparent to the user
- **Empty process**—A process that holds no active components associated with any application

If system resources must be reclaimed, Android will first terminate empty processes, followed by background processes, and so forth. Processes are assigned an importance ranking, and Android attempts to assign a process as high a ranking as possible. For example, if a process is providing a service and is also visible, it will be assigned the more-important visible classification.

Furthermore, Android development practices suggest following the guidelines of the process life cycle. When these guidelines are followed, the state of a process will be saved prior to termination and resumed at its saved state if the user navigates back to the application.

**PARTICIPATION
ACTIVITY**

2.3.3: Section review questions.



1) What best describes the return value from `fork()` for the child process?



- ☐ value < 0
- ☐ value of 0
- ☐ value > 0

2) What best describes what occurs when a child process calls the **exec ()** system call?



- ☐ The parent must wait for the child to terminate.
- ☐ The child is a duplicate of the parent.
- ☐ The child process has a new program loaded into it.

3) What best describes the return value from **fork ()** to the parent process after the parent successfully creates a child process?



- ☐ value < 0
- ☐ value of 0
- ☐ value > 0

4) A UNIX/Linux process that calls **exit ()** to terminate, but whose parent has not yet called **wait ()**, is known as a _____ process.



- ☐ zombie
- ☐ orphan
- ☐ terminated

5) A UNIX/Linux _____ process is assigned as the parent to orphan processes.



- ☐ zombie
- ☐ systemd
- ☐ sshd

Section glossary

tree: A data structure that can be used to represent data hierarchically; data values in a tree structure are linked through parent-child relationships.

process identifier (pid): A unique value for each process in the system that can be used as an index to access various attributes of a process within the kernel.

cascading termination: A technique in which, when a process is ended, all of its children are ended as well.

zombie: A process that has terminated but whose parent has not yet called `wait()` to collect its state and accounting information.

orphan: The child of a parent process that terminates in a system that does not require a terminating parent to cause its children to be terminated.

2.4 Interprocess communication


Processes executing concurrently in the operating system may be either independent processes or cooperating processes. A process is **independent** if it does not share data with any other processes executing in the system. A process is **cooperating** if it can affect or be affected by the other processes executing in the system. Clearly, any process that shares data with other processes is a cooperating process.

There are several reasons for providing an environment that allows process cooperation:

- **Information sharing.** Since several applications may be interested in the same piece of information (for instance, copying and pasting), we must provide an environment to allow concurrent access to such information.
- **Computation speedup.** If we want a particular task to run faster, we must break it into subtasks, each of which will be executing in parallel with the others. Notice that such a speedup can be achieved only if the computer has multiple processing cores.
- **Modularity.** We may want to construct the system in a modular fashion, dividing the system functions into separate processes or threads, as we discussed in the previous chapter Operating-System Structures.

Cooperating processes require an **interprocess communication** (IPC) mechanism that will allow them to exchange data— that is, send data to and receive data from each other. There are two fundamental models of interprocess communication: **shared memory** and **message passing**. In the shared-memory model, a region of memory that is shared by the cooperating processes is established. Processes can then exchange information by reading and writing data to the shared region. In the message-passing model, communication takes place by means of messages exchanged between the cooperating processes. The two communications models are contrasted in Figure [2.4.1](#).

Figure 2.4.1: Communications models. (a) Shared memory. (b) Message passing.



Multiprocess architecture—chrome browser

Many websites contain active content, such as JavaScript and HTML5, to provide a rich and dynamic web-browsing experience. Unfortunately, these web applications may also contain software bugs, which can result in sluggish response times and can even cause the web browser to crash. This isn't a big problem in a web browser that displays content from only one website. But most contemporary web browsers provide tabbed browsing, which allows a single instance of a web browser application to open several websites at the same time, with each site in a separate tab. To switch between the different sites, a user need only click on the appropriate tab. This arrangement is illustrated below:



A problem with this approach is that if a web application in any tab crashes, the entire process—including all other tabs displaying additional websites—crashes as well.

Google's Chrome web browser was designed to address this issue by using a multiprocess architecture. Chrome identifies three different types of processes: browser, renderers, and plug-ins.

- The **browser** process is responsible for managing the user interface as well as disk and network I/O. A new browser process is created when Chrome is started. Only one browser process is created.

- **Renderer** processes contain logic for rendering web pages. Thus, they contain the logic for handling HTML, Javascript, images, and so forth. As a general rule, a new renderer process is created for each website opened in a new tab, and so several renderer processes may be active at the same time.
- A **plug-in** process is created for each type of plug-in or Chrome extension in use. Plug-in processes contain the code for the plug-in as well as additional code that enables the plug-in to communicate with associated renderer processes and the browser process.

The advantage of the multiprocess approach is that websites run in isolation from one another. If one website crashes, only its renderer process is affected; all other processes remain unharmed. Furthermore, renderer processes run in a **sandbox**, which means that access to disk and network I/O is restricted, minimizing the effects of any security exploits.

Both of the models just mentioned are common in operating systems, and many systems implement both. Message passing is useful for exchanging smaller amounts of data, because no conflicts need be avoided. Message passing is also easier to implement in a distributed system than shared memory. (Although there are systems that provide distributed shared memory, we do not consider them in this text.) Shared memory can be faster than message passing, since message-passing systems are typically implemented using system calls and thus require the more time-consuming task of kernel intervention. In shared-memory systems, system calls are required only to establish shared-memory regions. Once shared memory is established, all accesses are treated as routine memory accesses, and no assistance from the kernel is required.

In Section [Missing Reference : IPC in shared-memory systems](#) and Section [Missing Reference : IPC in message-passing systems](#) we explore shared-memory and message-passing systems in more detail.

**PARTICIPATION
ACTIVITY**

2.4.1: Section review question.



What type of process is created by the Chrome browser for each tab that is opened?



Check

[Show answer](#)

**PARTICIPATION
ACTIVITY**

2.4.2: Section review question.



What best explains why shared memory is typically faster than message passing?



- ☐ Message passing is implemented in system calls which requires kernel intervention; shared memory only requires a system call to establish the shared memory segment.
- ☐ Message passing is typically used for sharing large amounts of data.
- ☐ Message passing is only used in distributed systems where communication takes place across a slower network.

Section glossary

interprocess communication (IPC): Communication between processes.

shared memory: In interprocess communication, a section of memory shared by multiple processes and used for message passing.

message passing: In interprocess communication, a method of sharing data in which messages are sent and received by processes. Packets of information in predefined formats are moved between processes or between computers.

browser: A process that accepts input in the form of a URL (Uniform Resource Locator), or web address, and displays its contents on a screen.

renderer: A process that contains logic for rendering contents (such as web pages) onto a display.

plug-in: An add-on functionality that expands the primary functionality of a process (e.g., a web browser plug-in that displays a type of content different from what the browser can natively handle).

sandbox: A contained environment (e.g., a virtual machine).

2.5 The process concept

The process concept


A **process** is an instance of a program being executed by an OS. Ex: When a user opens a new application like a web browser or text editor, the OS creates a new process.

The OS itself is organized as a collection of processes.

The OS keeps track of each process using a **process control block (PCB)**: A data structure that holds information for a process, including the current instruction address, the execution stack, the set of resources used by the process, and the program being executed. The PCB is the concrete representation of a process.

PARTICIPATION ACTIVITY

2.5.1: OS creates a new PCB for every new process.



Animation content:

Static figure: A graphical representation of operating system processes and their corresponding process control blocks (PCBs). A display screen on the left with 3 icons labeled Edge, Skype, and Word, and a series of boxes on the right. Each PCB points to boxes labeled as Program and Resources. Additionally, horizontal arrows represent the program counters.

Step 1: Before a user logs in, only OS processes are running. Each process is represented by a PCB, which points to a program and the resources controlled by the process. Horizontal arrows represent program counters.

The display screen shows 2 lines: "ID: ____" and "Password: ____" to be completed by the user. Two OS processes appear, each consisting of a PCB box pointing a Program and a Data box.

Step 2: When User1 logs in, a new PCB is created. The PCB points to User1's program, which typically executes the graphical user interface, and to initial resources. The program counter points to the start of the program.

The screen with fields displaying ID: User1 and a masked password field. A new PCB with program and data labeled User1 appears next to the OS processes.

Step 3: User1 is running. The screen shows the icons of different applications.

The user's display screen now shows application icons labeled Edge, Skype, and Word, with a cursor arrow pointing at the Edge icon.

Step 4: When the user clicks the Microsoft Edge icon, a new PCB for Microsoft Edge is created by the OS.

Alongside the existing PCBs for OS processes and User1, a new PCB with program and data labeled Edge appears.

Animation captions:

1. Before a user logs in, only OS processes are running. Each process is represented by a PCB, which points to a program and the resources controlled by the process. Horizontal arrows represent program counters.
2. When User1 logs in, a new PCB is created. The PCB points to User1's program, which typically executes the graphical user interface, and to initial resources. The program counter points to the start of the program.
3. User1 is running. The screen shows the icons of different applications.
4. When the user clicks the Microsoft Edge icon, a new PCB for Microsoft Edge is created by the OS.

PARTICIPATION ACTIVITY

2.5.2: Basic principles of a process.



1) An OS uses a PCB to represent a process.



- ☐ True
- ☐ False

2) The PCB is created by the process when execution starts.



- ☐ True
- ☐ False

3) The PCB becomes part of the program being executed by a process.

- ☐ True
- ☐ False

4) Two processes can be executing the same program.



- ☐ True
- ☐ False

The process concept by analogy

A good way to understand the concept of a process is to consider an analogy from real life.

PARTICIPATION ACTIVITY

2.5.3: Correspondence between the process concept and the employment agency analogy.



Animation content:

Static figure: An Employment agency comparison to OS. The Employment agency consists of several boxes. "Resources" points to "Jobs status record", which in turn points to "Job". A "Worker" icon points to step 1 in Job. The OS has boxes parallel to those in the Employment agency. "Resources" points to "PCB", which in turn points to "Program". A "CPU" icon points to instr 1 in Program.

Step 1: An agency employs a worker to perform various jobs. By analogy, the OS uses the CPU to execute different programs. "Worker" appears in Employment agency and "CPU" appears in OS.

Step 2: The worker performs jobs consisting of multiple steps, analogous to the CPU, which executes programs consisting of multiple instructions.

"Job" appears in Employment agency and "Program" appears in OS.

Step 3: The agency sends the worker to a location where the job takes place. The work location corresponds to the memory where the program execution takes place.

Step 4: The agency creates a job status record, which is the dynamic representation of the job's progress, analogous to the PCB, which is the representation of the process.

"Job status record" pointing to "Job" appears in Employment agency and "PCB" pointing to "Program" appears in OS.

Step 5: The job status record also keeps track of the resources assigned to the worker in the same way as the PCB keeps track of resources assigned to the process.

"Resources" pointing to "Job status record" in Employment agency and to "PCB" in OS appear.

Animation captions:

1. An agency employs a worker to perform various jobs. By analogy, the OS uses the CPU to execute different programs.
2. The worker performs jobs consisting of multiple steps, analogous to the CPU, which executes programs consisting of multiple instructions.
3. The agency sends the worker to a location where the job takes place. The work location corresponds to the memory where the program execution takes place.
4. The agency creates a job status record, which is the dynamic representation of the job's progress, analogous to the PCB, which is the representation of the process.
5. The job status record also keeps track of the resources assigned to the worker in the same way as the PCB keeps track of resources assigned to the process.

PARTICIPATION ACTIVITY

2.5.4: Understanding the process concept via a student library analogy.



A process can be compared to the act of reading a book checked out by a student from the reserved books section of a library.

How to use this tool ▼

memory

OS

PCB

program

CPU


The student corresponds to the ____.

	The book corresponds to the ____.
	The library reading room corresponds to the ____.
	The librarian corresponds to the ____.
	The check-out record corresponds to the ____.

Reset

PARTICIPATION
ACTIVITY

2.5.5: Process states and transitions.



Animation content:

Static figure: A state diagram with three states that a process can be in: Ready, Running, and Blocked. Arrows indicate possible transitions between states with labels describing the cause of the transition: Request resource points from Running to Blocked; Release resource points from Blocked to Ready; there are two bidirectional arrows between Running and Ready labeled OS scheduling, implying that a process can move back and forth between these states.

Step 1: A newly created process starts in the ready state. Ready state is highlighted.

Step 2: When the OS selects the process for execution, the process's state changes to running. Running state is highlighted

Step 3: When the OS interrupts the process, for example to let some other process use the CPU, the current process's state changes back to ready. Ready state is highlighted.

Step 4: When the OS again restarts the process, the process's state changes back to running. Running state is highlighted

Step 5: When the process requests a currently unavailable resource, such as a file already open by another user or some input data, the process's state changes to blocked. Blocked state is highlighted

Step 6: When some other process releases the resource or generates the data needed by the blocked process, the OS moves the blocked process back to ready. Ready state is highlighted.

Animation captions:

1. A newly created process starts in the ready state.
2. When the OS selects the process for execution, the process's state changes to running.
3. When the OS interrupts the process, for example to let some other process use the CPU, the current process's state changes back to ready.
4. When the OS again restarts the process, the process's state changes back to running.
5. When the process requests a currently unavailable resource, such as a file already open by another user or some input data, the process's state changes to blocked.
6. When some other process releases the resource or generates the data needed by the blocked process, the OS moves the blocked process back to ready.

PARTICIPATION ACTIVITY

2.5.6: Process states and transitions.



1) The transition (ready -> running) of a process p is caused by ____.



- ☐ the OS
- ☐ the process p itself
- ☐ some other process

2) The transition (running -> blocked) of a process p is caused by ____.



- ☐ the OS
- ☐ the process p itself
- ☐ some other process

3) The transition (running -> ready) of a process p is caused by ____.

- ☐ the OS
- ☐ the process p itself
- ☐ some other process

4) The transition (blocked -> ready) of a process p is caused by ____.

- ☐ the OS
- ☐ the process p itself
- ☐ some other process


Additional process states and transitions

Some OSs provide a more complex set of states and transitions. The additional states may include the following:

- A newly created process is placed into the **new state** before the process is allowed to compete for the CPU. Ex: The OS may want to regulate the number of processes competing for the CPU.
- A process is placed into the **terminated state** when execution can no longer continue but before the PCB is deleted. Ex: The OS may want to examine the final state of a process that committed a fatal error.
- A process may be placed into the **suspended state** even though the CPU and all resources are available. Ex: The OS may want to stop a process to allow debugging or to regulate performance.

PARTICIPATION ACTIVITY

2.5.7: Additional state transitions.



Animation content:

Static figure: A state diagram with six states that a process can be in: New, Ready, Running, Blocked, Suspended, and Terminated. Arrows labeled with different transition names connect these states, indicating possible state changes. The transitions include Activate from New to Ready, OS scheduling between Ready to Running, Request resource from Running to Blocked, Release resource from Blocked to Ready and, Suspend from Running or Ready or Blocked to Suspended, and End from Running to Terminated. The Suspended state has additional transitions labeled Reactivate, which lead to Ready or Blocked, depending on the state from which the process was suspended.

Additional details:

Step 1: A newly created process starts in the new state. New is highlighted.

Step 2: The OS moves the process from new to ready when the process is allowed to compete for the CPU. The transition is called Activate. Ready is highlighted.

Step 3: From ready, the process can move repeatedly between running, blocked, and ready as before. Ready, Running, and Blocked are highlighted.

Step 4: In addition, the OS can move the process from the running, ready, or blocked state to a suspended state to temporarily prevent the process from making any progress. Suspended is highlighted.

Step 5: If the process was suspended in the blocked state, then upon reactivation the process returns to the blocked state. Blocked is highlighted.

Step 6: If the suspension occurred in the running or ready state, then upon reactivation the process is returned to the ready state. Suspended is highlighted, then Ready.

Step 7: When a running process finishes all work or causes a fatal error, the process's state changes to terminated. The OS can also perform the transition to terminated from any other state. Running is highlighted, then Terminated.

Animation captions:

1. A newly created process starts in the new state.
2. The OS moves the process from new to ready when the process is allowed to compete for the CPU. The transition is called "Activate".
3. From ready, the process can move repeatedly between running, blocked, and ready as before.
4. In addition, the OS can move the process from the running, ready, or blocked state to a suspended state to temporarily prevent the process from making any progress.
5. If the process was suspended in the blocked state, then upon reactivation the process returns to the blocked state.

6. If the suspension occurred in the running or ready state, then upon reactivation the process is returned to the ready state.
7. When a running process finishes all work or causes a fatal error, the process's state changes to terminated. The OS can also perform the transition to terminated from any other state.

**PARTICIPATION
ACTIVITY**

2.5.8: Additional process states and transitions.



Copyright © 2023, Addison-Wesley, a division of Pearson Education, Inc. All rights reserved.

- 1) When a newly created process p is ready to compete for the CPU, p moves itself from the new to the ready state.
☐ True
☐ False
- 2) When a ready process p wishes to temporarily stop competing for the CPU, p moves itself to the suspended state.
☐ True
☐ False
- 3) The transition of a process p from running to terminated can be caused by p itself.
☐ True
☐ False
- 4) When a blocked process p is suspended and the resource needed by p becomes available, p is moved to the ready state.
☐ True
☐ False



**PARTICIPATION
ACTIVITY**

2.5.9: Process reactivation.



Copyright © 2023, Addison-Wesley, a division of Pearson Education, Inc. All rights reserved.

When a process p is moved from the current state (running, ready, or blocked) to the suspended state, then upon reactivation,

- ☐ p must be moved back to the same state where the suspension occurred
- ☐ p can always safely be moved to the ready state
- ☐ p is moved to either the ready or the blocked state

The context switch


The CPU is always running one process at a time. Multiple processes can share a single CPU by taking turns. A **context switch** is the transfer of control from one process to another and may be triggered by the OS for scheduling reasons or caused by the inability of the current process to continue.

Each time one process stops executing to allow another one to resume, the OS must save all information about the stopped process. This information is restored when the process again gets a chance to run. The temporary interruption is fully transparent to the process.

The information that needs to be saved is called the CPU state. The **CPU state** consists of all intermediate values held in any CPU registers and hardware flags at the time of the interruption. One of the registers is the program counter, which determines the next instruction to execute after the process is restored.

PARTICIPATION ACTIVITY

2.5.10: Context switches between processes p and q.



Animation content:

Static figure: A representation of two processes, p and q, and their interaction with the CPU and process control blocks (PCBs). Column 1: two vertical timelines with arrows between them indicating the direction of interactions between p and q. Column 2: A CPU, which contains the current state of a process. Column 3: Two PCBs, each containing the name p or q, and the Saved CPU state.

Step 1: Process p is running. The current state of p is in the CPU, while p's PCB contains an out-of-date copy.

A down-arrow from p is highlighted to indicate that process p is the one currently running. The PCB for process p visually represents a copy of the CPU state, which is currently not up to date.

Step 2: A hardware interrupt transfers control to the OS, which saves the current CPU state in p's PCB. A horizontal arrow from p to q indicates that p stopped and q continues to run.

An arrow from the CPU pointing towards process p's PCB, labeled Save, indicates that the state of the CPU is being copied to the PCB of process p.

Step 3: If the OS chooses process q to continue, the current CPU state is overwritten with q's last state from q's PCB.

An arrow labeled Restore points from process q's PCB to the CPU, indicating that the saved state of q is copied to the CPU.

Step 4: As q continues running, the saved state in q's PCB gets out of date. q's current state is maintained in the CPU.

A down-arrow in the q timeline is highlighted to indicate that process q is the one currently running. The saved state in q's PCB fades.

Step 5: Another interrupt causes the current CPU state to be saved in q's PCB. The OS then may restart p by restoring p's saved state in the CPU.

A horizontal arrow from q to p indicates that q stopped and p continues to run.

An arrow from CPU state to q's PCB indicates the action of saving the CPU's present state into process q's PCB. Another arrow labeled Restore points from process p's PCB to the CPU to indicate that process p's state is being restored.

Step 6: As p continues running, the saved state in the PCB gets out of date.

A down-arrow in the p timeline is highlighted to indicate that process p is the one currently running. The saved state in p's PCB fades.

Animation captions:

1. Process p is running. The current state of p is in the CPU, while p's PCB contains an out-of-date copy.
2. A hardware interrupt transfers control to the OS, which saves the current CPU state in p's PCB.

3. If the OS chooses process q to continue, the current CPU state is overwritten with q's last state from q's PCB.
4. As q continues running, the saved state in q's PCB gets out of date. q's current state is maintained in the CPU.
5. Another interrupt causes the current CPU state to be saved in q's PCB. The OS then may restart p by restoring p's saved state in the CPU.
6. As p continues running, the saved state in the PCB gets out of date.

**PARTICIPATION
ACTIVITY**

2.5.11: Who causes a context switch.

CS 510 10435 202617-1



A context switch for a currently running process p may be caused by ____.

1) process p itself



- ☐ True
- ☐ False

2) the OS



- ☐ True
- ☐ False

3) some other process q



- ☐ True
- ☐ False

**PARTICIPATION
ACTIVITY**

2.5.12: Current CPU state versus copy of CPU state in PCB.



The PCB contains an up-to-date copy of the CPU state of a process p ____ .



- ☐ when p's state is running
- ☐ when p's state is ready or blocked
- ☐ at all times

**CHALLENGE
ACTIVITY**

2.5.1: State transitions.



726670.4301284.qx3zqy7

Start

Which of the following state transitions are possible?

- ☐ Ready -> blocked
- ☐ Suspended -> running
- ☐ Running -> suspended
- ☐ Running -> ready
- ☐ Blocked -> running
- ☐ Blocked -> ready

1	2	3	4	5
---	---	---	---	---

Check

Next



EXERCISE

2.5.1: State transitions.



Indicate whether each series of state transitions is valid or invalid.

- (a) new → ready → suspended → blocked → suspended.
- (b) suspended → blocked → suspended → blocked → ready → running.
- (c) running → blocked → ready → blocked → suspended.
- (d) new → ready → running → ready → new.