

7.1 Background

Introduction

In the chapter Main Memory, we discussed various memory-management strategies used in computer systems. All these strategies have the same goal: to keep many processes in memory simultaneously to allow multiprogramming. However, they tend to require that an entire process be in memory before it can execute.

Virtual memory is a technique that allows the execution of processes that are not completely in memory. One major advantage of this scheme is that programs can be larger than physical memory. Further, virtual memory abstracts main memory into an extremely large, uniform array of storage, separating logical memory as viewed by the programmer from physical memory. This technique frees programmers from the concerns of memory-storage limitations. Virtual memory also allows processes to share files and libraries, and to implement shared memory. In addition, it provides an efficient mechanism for process creation. Virtual memory is not easy to implement, however, and may substantially decrease performance if it is used carelessly. In this chapter, we provide a detailed overview of virtual memory, examine how it is implemented, and explore its complexity and benefits.

Chapter objectives

- Define virtual memory and describe its benefits.
 - Illustrate how pages are loaded into memory using demand paging.
 - Apply the FIFO, optimal, and LRU page-replacement algorithms.
 - Describe the working set of a process, and explain how it is related to program locality.
 - Describe how Linux, Windows 10, and Solaris manage virtual memory.
 - Design a virtual memory manager simulation in the C programming language.
-

Background

The memory-management algorithms outlined in the chapter Main Memory are necessary because of one basic requirement: the instructions being executed must be in physical memory. The first approach to meeting this requirement is to place the entire logical address space in physical memory. Dynamic linking can help to ease this restriction, but it generally requires special precautions and extra work by the programmer.

The requirement that instructions must be in physical memory to be executed seems both necessary and reasonable; but it is also unfortunate, since it limits the size of a program to the size of physical memory. In fact, an examination of real programs shows us that, in many cases, the entire program is not needed. For instance, consider the following:

- Programs often have code to handle unusual error conditions. Since these errors seldom, if ever, occur in practice, this code is almost never executed.
- Arrays, lists, and tables are often allocated more memory than they actually need. An array may be declared 100 by 100 elements, even though it is seldom larger than 10 by 10 elements.
- Certain options and features of a program may be used rarely. For instance, the routines on U.S. government computers that balance the budget have not been used in many years.

Even in those cases where the entire program is needed, it may not all be needed at the same time.


The ability to execute a program that is only partially in memory would confer many benefits:

- A program would no longer be constrained by the amount of physical memory that is available. Users would be able to write programs for an extremely large **virtual** address space, simplifying the programming task.
- Because each program could take less physical memory, more programs could be run at the same time, with a corresponding increase in CPU utilization and throughput but with no increase in response time or turnaround time.
- Less I/O would be needed to load or swap portions of programs into memory, so each program would run faster.

Thus, running a program that is not entirely in memory would benefit both the system and its users.


Virtual memory involves the separation of logical memory as perceived by developers from physical memory. This separation allows an extremely large virtual memory to be provided for programmers when only a smaller physical memory is available (Figure [7.1.1](#)). Virtual memory makes the task of programming much easier, because the programmer no longer needs to worry about the amount of physical memory available; she can concentrate instead on programming the problem that is to be solved.

Figure 7.1.1: Diagram showing virtual memory that is larger than physical memory.



The **virtual address space** of a process refers to the logical (or virtual) view of how a process is stored in memory. Typically, this view is that a process begins at a certain logical address—say, address 0—and exists in contiguous memory, as shown in Figure 7.1.2. Recall from the chapter Main Memory, though, that in fact physical memory is organized in page frames and that the physical page frames assigned to a process may not be contiguous. It is up to the memory-management unit (MMU) to map logical pages to physical page frames in memory.

Figure 7.1.2: Virtual address space of a process in memory.




Note in Figure [7.1.2](#) that we allow the heap to grow upward in memory as it is used for dynamic memory allocation. Similarly, we allow for the stack to grow downward in memory through successive function calls. The large blank space (or hole) between the heap and the stack is part of the virtual address space but will require actual physical pages only if the heap or stack grows. Virtual address spaces that include holes are known as **sparse** address spaces. Using a sparse address space is beneficial because the holes can be filled as the stack or heap segments grow or if we wish to dynamically link libraries (or possibly other shared objects) during program execution.

In addition to separating logical memory from physical memory, virtual memory allows files and memory to be shared by two or more processes through page sharing (Section Shared pages). This leads to the following benefits:

- System libraries such as the standard C library can be shared by several processes through mapping of the shared object into a virtual address space. Although each process considers the libraries to be part of its virtual address space, the actual pages where the libraries reside in physical memory are shared by all the processes (Figure [7.1.3](#)). Typically, a library is mapped read-only into the space of each process that is linked with it.
- Similarly, processes can share memory. Recall from the chapter Processes that two or more processes can communicate through the use of shared memory. Virtual memory allows one process to create a region of memory that it can share with another process. Processes sharing this region consider it part of their virtual address space, yet the actual physical pages of memory are shared, much as is illustrated in Figure [7.1.3](#).
- Pages can be shared during process creation with the `fork()` system call, thus speeding up process creation.

Figure 7.1.3: Shared library using virtual memory.



We further explore these—and other—benefits of virtual memory later in this chapter. First, though, we discuss implementing virtual memory through demand paging.

**PARTICIPATION
ACTIVITY**

7.1.1: Section review questions.

1) Virtual memory _____.

- ☐ assumes a process resides entirely in physical memory
- ☐ allows only portions of a process be in physical memory
- ☐ limits the amount of physical memory allocated to a process

2) In systems that support virtual memory, _____.

- ☐ virtual memory is separated from logical memory.
- ☐ physical memory is separated from secondary storage.
- ☐ physical memory is separated from logical memory.

3) In general, using virtual memory technology decreases the degree of multiprogramming in a system.

- ☐ True
- ☐ False



Section glossary

virtual memory: A technique that allows the execution of a process that is not completely in memory. Also, separation of computer memory address space from physical into logical, allowing easier programming and larger name space.

virtual address space: The logical view of how a process is stored in memory.

sparse: In memory management, a term describing a page table that has noncontiguous, scattered entries. A sparse address space has many holes.

7.2 Demand paging

Consider how an executable program might be loaded from secondary storage into memory. One option is to load the entire program in physical memory at program execution time. However, a problem with this approach is that we may not initially **need** the entire program in memory. Suppose a program starts with a list of available options from which the user is to select. Loading the entire program into memory results in loading the executable code for **all** options, regardless of whether or not an option is ultimately selected by the user.


An alternative strategy is to load pages only as they are needed. This technique is known as **demand paging** and is commonly used in virtual memory systems. With demand-paged virtual memory, pages are loaded only when they are **demande**d during program execution. Pages that are never accessed are thus never loaded into physical memory. A demand-paging system is similar to a paging system with swapping (Section Swapping with paging) where processes reside in secondary memory (usually an HDD or NVM device). Demand paging explains one of the primary benefits of virtual memory—by loading only the portions of programs that are needed, memory is used more efficiently.

Basic concepts

The general concept behind demand paging, as mentioned, is to load a page in memory only when it is needed. As a result, while a process is executing, some pages will be in memory, and some will be in secondary storage. Thus, we need some form of hardware support to distinguish between the two. The

valid-invalid bit scheme described in Section Protection can be used for this purpose. This time, however, when the bit is set to "valid," the associated page is both legal and in memory. If the bit is set to "invalid," the page either is not valid (that is, not in the logical address space of the process) or is valid but is currently in secondary storage. The page-table entry for a page that is brought into memory is set as usual, but the page-table entry for a page that is not currently in memory is simply marked invalid. This situation is depicted in Figure 7.2.1. (Notice that marking a page invalid will have no effect if the process never attempts to access that page.)

Figure 7.2.1: Page table when some pages are not in main memory.




But what happens if the process tries to access a page that was not brought into memory? Access to a page marked invalid causes a **page fault**. The paging hardware, in translating the address through the page table, will notice that the invalid bit is set, causing a trap to the operating system. This trap is the result of the operating system's failure to bring the desired page into memory. The procedure for handling this page fault is straightforward (animation below):

1. We check an internal table (usually kept with the process control block) for this process to determine whether the reference was a valid or an invalid memory access.
2. If the reference was invalid, we terminate the process. If it was valid but we have not yet brought in that page, we now page it in.
3. We find a free frame (by taking one from the free-frame list, for example).
4. We schedule a secondary storage operation to read the desired page into the newly allocated frame.
5. When the storage read is complete, we modify the internal table kept with the process and the page table to indicate that the page is now in memory.
6. We restart the instruction that was interrupted by the trap. The process can now access the page as though it had always been in memory.

**PARTICIPATION
ACTIVITY**

7.2.1: Steps in handling a page fault.



Animation content:

Step 1: The instruction "load M" executes, accessing the page table for the address containing "M".
Step 2: The page is marked as invalid, causing a page fault trap to the kernel. The operating system checks if the reference is legal or not. If it is illegal, the process is terminated. In this case it is valid, but the page is not currently in a frame in physical memory. Step 3: The kernel finds a free frame in physical memory (taking one from the free frame list for example) and schedules a secondary storage operation to read the desired page into the newly allocated frame. Step 4: The frame number is loaded

into the page table entry. Step 5: We restart the instruction that was interrupted by the trap - the process can now access the page as though it had always been in memory.

Animation captions:

1. The instruction "load M" executes, accessing the page table for the address containing "M".
2. The page is marked as invalid, causing a page fault trap to the kernel. The operating system checks if the reference is legal or not. If it is illegal, the process is terminated. In this case it is valid, but the page is not currently in a frame in physical memory.
3. The kernel finds a free frame in physical memory (taking one from the free frame list for example) and schedules a secondary storage operation to read the desired page into the newly allocated frame.
4. The frame number is loaded into the page table entry.
5. We restart the instruction that was interrupted by the trap - the process can now access the page as though it had always been in memory.

In the extreme case, we can start executing a process with **no** pages in memory. When the operating system sets the instruction pointer to the first instruction of the process, which is on a non-memory-resident page, the process immediately faults for the page. After this page is brought into memory, the process continues to execute, faulting as necessary until every page that it needs is in memory. At that point, it can execute with no more faults. This scheme is **pure demand paging**: never bring a page into memory until it is required.

Theoretically, some programs could access several new pages of memory with each instruction execution (one page for the instruction and many for data), possibly causing multiple page faults per instruction. This situation would result in unacceptable system performance. Fortunately, analysis of running processes shows that this behavior is exceedingly unlikely. Programs tend to have **locality of reference**, described in Section Cause of thrashing, which results in reasonable performance from demand paging.

The hardware to support demand paging is the same as the hardware for paging and swapping:

- **Page table.** This table has the ability to mark an entry invalid through a valid-invalid bit or a special value of protection bits.
- **Secondary memory.** This memory holds those pages that are not present in main memory. The secondary memory is usually a high-speed disk or NVM device. It is known as the swap device, and the section of storage used for this purpose is known as **swap space**. Swap-space allocation is discussed in the chapter Mass-Storage Structure.

A crucial requirement for demand paging is the ability to restart any instruction after a page fault. Because we save the state (registers, condition code, instruction counter) of the interrupted process when the page fault occurs, we must be able to restart the process in **exactly** the same place and state, except that the desired page is now in memory and is accessible. In most cases, this requirement is easy to meet. A page fault may occur at any memory reference. If the page fault occurs on the instruction fetch, we can restart by fetching the instruction again. If a page fault occurs while we are fetching an operand, we must fetch and decode the instruction again and then fetch the operand.

As a worst-case example, consider a three-address instruction such as ADD the content of A to B, placing the result in C. These are the steps to execute this instruction:

1. Fetch and decode the instruction (ADD).
2. Fetch A.
3. Fetch B.
4. Add A and B.
5. Store the sum in C.

If we fault when we try to store in C (because C is in a page not currently in memory), we will have to get the desired page, bring it in, correct the page table, and restart the instruction. The restart will require fetching the instruction again, decoding it again, fetching the two operands again, and then adding again. However, there is not much repeated work (less than one complete instruction), and the repetition is necessary only when a page fault occurs.

The major difficulty arises when one instruction may modify several different locations. For example, consider the IBM System 360/370 MVC (move character) instruction, which can move up to 256 bytes from one location to another (possibly overlapping) location. If either block (source or destination) straddles a page boundary, a page fault might occur after the move is partially done. In addition, if the source and destination blocks overlap, the source block may have been modified, in which case we cannot simply restart the instruction.

This problem can be solved in two different ways. In one solution, the microcode computes and attempts to access both ends of both blocks. If a page fault is going to occur, it will happen at this step, before anything is modified. The move can then take place; we know that no page fault can occur, since all the relevant pages are in memory. The other solution uses temporary registers to hold the values of overwritten locations. If there is a page fault, all the old values are written back into memory before the trap occurs. This action restores memory to its state before the instruction was started, so that the instruction can be repeated.

This is by no means the only architectural problem resulting from adding paging to an existing architecture to allow demand paging, but it illustrates some of the difficulties involved. Paging is added between the CPU and the memory in a computer system. It should be entirely transparent to a process. Thus, people often assume that paging can be added to any system. Although this assumption is true for a non-demand-paging environment, where a page fault represents a fatal error, it is not true where a page fault means only that an additional page must be brought into memory and the process restarted.

Free-frame list

When a page fault occurs, the operating system must bring the desired page from secondary storage into main memory. To resolve page faults, most operating systems maintain a **free-frame list**, a pool of free frames for satisfying such requests (Figure 7.2.2). (Free frames must also be allocated when the stack or heap segments from a process expand.) Operating systems typically allocate free frames using a technique known as **zero-fill-on-demand**. Zero-fill-on-demand frames are "zeroed-out" before being allocated, thus erasing their previous contents. (Consider the potential security implications of **not** clearing out the contents of a frame before reassigning it.)

Figure 7.2.2: List of free frames.



When a system starts up, all available memory is placed on the free-frame list. As free frames are requested (for example, through demand paging), the size of the free-frame list shrinks. At some point, the list either falls to zero or falls below a certain threshold, at which point it must be repopulated. We cover strategies for both of these situations in Section [7.4](#).

Performance of demand paging

Demand paging can significantly affect the performance of a computer system. To see why, let's compute the **effective access time** for a demand-paged memory. Assume the memory-access time, denoted *ma*, is 10 nanoseconds. As long as we have no page faults, the effective access time is equal to the memory access time. If, however, a page fault occurs, we must first read the relevant page from secondary storage and then access the desired word.

Let *p* be the probability of a page fault ($0 \leq p \leq 1$). We would expect *p* to be close to zero—that is, we would expect to have only a few page faults. The effective access time is then

$$\text{effective access time} = (1 - p) \times ma + p \times \text{page fault time}.$$

To compute the effective access time, we must know how much time is needed to service a page fault. A page fault causes the following sequence to occur:

1. Trap to the operating system.
2. Save the registers and process state.
3. Determine that the interrupt was a page fault.
4. Check that the page reference was legal, and determine the location of the page in secondary storage.
5. Issue a read from the storage to a free frame:
 - a. Wait in a queue until the read request is serviced.
 - b. Wait for the device seek and/or latency time.
 - c. Begin the transfer of the page to a free frame.
6. While waiting, allocate the CPU core to some other process.
7. Receive an interrupt from the storage I/O subsystem (I/O completed).
8. Save the registers and process state for the other process (if step 6 is executed).
9. Determine that the interrupt was from the secondary storage device.

10. Correct the page table and other tables to show that the desired page is now in memory.
11. Wait for the CPU core to be allocated to this process again.
12. Restore the registers, process state, and new page table, and then resume the interrupted instruction.

Not all of these steps are necessary in every case. For example, we are assuming that, in step 6, the CPU is allocated to another process while the I/O occurs. This arrangement allows multiprogramming to maintain CPU utilization but requires additional time to resume the page-fault service routine when the I/O transfer is complete.

In any case, there are three major task components of the page-fault service time:

1. Service the page-fault interrupt.
2. Read in the page.
3. Restart the process.

The first and third tasks can be reduced, with careful coding, to several hundred instructions. These tasks may take from 1 to 100 microseconds each. Let's consider the case of HDDs being used as the paging device. The page-switch time will probably be close to 8 milliseconds. (A typical hard disk has an average latency of 3 milliseconds, a seek of 5 milliseconds, and a transfer time of 0.05 milliseconds. Thus, the total paging time is about 8 milliseconds, including hardware and software time.) Remember also that we are looking at only the device-service time. If a queue of processes is waiting for the device, we have to add queuing time as we wait for the paging device to be free to service our request, increasing even more the time to page in.

With an average page-fault service time of 8 milliseconds and a memory-access time of 200 nanoseconds, the effective access time in nanoseconds is

$$\begin{aligned}
 \text{effective access time} &= (1 - p) \times (200) + p \text{ (8 milliseconds)} \\
 &= (1 - p) \times 200 + p \times 8,000,000 \\
 &= 200 + 7,999,800 \times p.
 \end{aligned}$$

We see, then, that the effective access time is directly proportional to the **page-fault rate**. If one access out of 1,000 causes a page fault, the effective access time is 8.2 microseconds. The computer will be slowed down by a factor of 40 because of demand paging! If we want performance degradation to be less than 10 percent, we need to keep the probability of page faults at the following level:

$$\begin{aligned}
 220 &> 200 + 7,999,800 \times p, \\
 20 &> 7,999,800 \times p, \\
 p &< 0.0000025.
 \end{aligned}$$

That is, to keep the slowdown due to paging at a reasonable level, we can allow fewer than one memory access out of 399,990 to page-fault. In sum, it is important to keep the page-fault rate low in a demand-paging system. Otherwise, the effective access time increases, slowing process execution dramatically.

An additional aspect of demand paging is the handling and overall use of swap space. I/O to swap space is generally faster than that to the file system. It is faster because swap space is allocated in much larger blocks, and file lookups and indirect allocation methods are not used (Chapter Mass-Storage Structure). One option for the system to gain better paging throughput is by copying an entire file image into the swap space at process startup and then performing demand paging from the swap space. The obvious disadvantage of this approach is the copying of the file image at program start-up. A second option—and one practiced by several operating systems, including Linux and Windows—is to demand-page from the file system initially but to write the pages to swap space as they are replaced. This approach will ensure that only needed pages are read from the file system but that all subsequent paging is done from swap space.

Some systems attempt to limit the amount of swap space used through demand paging of binary executable files. Demand pages for such files are brought directly from the file system. However, when page replacement is called for, these frames can simply be overwritten (because they are never modified), and the pages can be read in from the file system again if needed. Using this approach, the file system itself serves as the backing store. However, swap space must still be used for pages not associated with a file (known as **anonymous memory**); these pages include the stack and heap for a process. This method appears to be a good compromise and is used in several systems, including Linux and BSD UNIX.

As described in Section Swapping on mobile systems, mobile operating systems typically do not support swapping. Instead, these systems demand-page from the file system and reclaim read-only pages (such as code) from applications if memory becomes constrained. Such data can be demand-paged from the file system if it is later needed. Under iOS, anonymous memory pages are never reclaimed from an application unless the application is terminated or explicitly releases the memory. In Section 7.7, we cover compressed memory, a commonly used alternative to swapping in mobile systems.

**PARTICIPATION
ACTIVITY**

7.2.2: Section review questions.



1) Referencing a page that is not yet loaded in main memory is known as _____ .

- ☐ demand paging
- ☐ a page fault
- ☐ a trap



2) Handling a page fault requires _____ .

- ☐ setting the valid/invalid bit to valid
- ☐ requires checking the value of the valid/invalid bit
- ☐ bringing the page from the backing store and storing it in a frame in physical memory



3) On a system with demand-paging, a process will experience a high page fault rate when the process begins execution.

- ☐ True
- ☐ False

4) If a memory access requires 100 ns and servicing a page fault requires 6 ms, what is the effective access time with page fault rate of $p = 0.25$?

- ☐ 4,500,025 ns
- ☐ 1,500,075 ns
- ☐ 1,500,025 ns

Section glossary

demand paging: In memory management, bringing in pages from storage as needed rather than, e.g., in their entirety at process load time.

page fault: A fault resulting from a reference to a non-memory-resident page of memory.

pure demand paging: A demand paging scheme wherein no page is brought into memory until it is referenced.

locality of reference: The tendency of processes to reference memory in patterns rather than randomly.

swap space: Secondary storage backing-store space used to store pages that are paged out of memory.

free-frame list: A kernel-maintained data structure containing the list of currently available free frames of physical memory.

zero-fill-on-demand: The writing of zeros into a page before it is made available to a process (to keep any old data from being available to the process).

effective access time: The measured or statistically calculated time it takes to access something; e.g., see effective memory-access time.

page-fault rate: A measure of how often a page fault occurs per memory access attempt.

anonymous memory: Memory not associated with a file. Pages not associated with a file, if dirty and paged out, must not lose their contents and are stored in swap space as anonymous memory.

7.3 Copy-on-write

In Section [7.2](#), we illustrated how a process can start quickly by demand-paging in the page containing the first instruction. However, process creation using the `fork()` system call may initially bypass the need for demand paging by using a technique similar to page sharing (covered in Section Shared pages). This technique provides rapid process creation and minimizes the number of new pages that must be allocated to the newly created process.

Recall that the `fork()` system call creates a child process that is a duplicate of its parent. Traditionally, `fork()` worked by creating a copy of the parent's address space for the child, duplicating the pages belonging to the parent. However, considering that many child processes invoke the `exec()` system call immediately after creation, the copying of the parent's address space may be unnecessary. Instead, we can use a technique known as **copy-on-write**, which works by allowing the parent and child processes initially to share the same pages. These shared pages are marked as copy-on-write pages, meaning that if either process writes to a shared page, a copy of the shared page is created. Copy-on-write is illustrated in

Figures 7.3.1 and 7.3.2, which show the contents of the physical memory before and after process 1 modifies page C.

Figure 7.3.1: Before process 1 modifies page C.





Figure 7.3.2: After process 1 modifies page C.



For example, assume that the child process attempts to modify a page containing portions of the stack, with the pages set to be copy-on-write. The operating system will obtain a frame from the free-frame list and create a copy of this page, mapping it to the address space of the child process. The child process will then modify its copied page and not the page belonging to the parent process. Obviously, when the copy-on-write technique is used, only the pages that are modified by either process are copied; all unmodified pages can be shared by the parent and child processes. Note, too, that only pages that can be modified need be marked as copy-on-write. Pages that cannot be modified (pages containing executable code) can be shared by the parent and child. Copy-on-write is a common technique used by several operating systems, including Windows, Linux, and macOS.

Several versions of UNIX (including Linux, macOS, and BSD UNIX) provide a variation of the `fork()` system call—`vfork()` (for **virtual memory fork**)—that operates differently from `fork()` with copy-on-write. With `vfork()`, the parent process is suspended, and the child process uses the address space of the parent. Because `vfork()` does not use copy-on-write, if the child process changes any pages of the parent's address space, the altered pages will be visible to the parent once it resumes. Therefore, `vfork()` must be used with caution to ensure that the child process does not modify the address space of the parent. `vfork()` is intended to be used when the child process calls `exec()` immediately after creation. Because no copying of pages takes place, `vfork()` is an extremely efficient method of process creation and is sometimes used to implement UNIX command-line shell interfaces.

**PARTICIPATION
ACTIVITY**

7.3.1: Section review questions.



_____ allows the parent and child processes to initially share the same pages, but when either process modifies a page, a copy of the shared page is created.



- ☐ copy-on-write
- ☐ zero-fill-on-demand
- ☐ `vfork()`

Section glossary

copy-on-write: Generally, the practice by which any write causes the data to first be copied and then modified, rather than overwritten. In virtual memory, on a write attempt to a shared page, the page is first copied, and the write is made to that copy.

virtual memory fork: The `vfork()` system call, which forks a child process, suspends the parent, and lets the child share the parent's address space for both read and write operations (changes are visible to the parent).

7.4 Page replacement

In our earlier discussion of the page-fault rate, we assumed that each page faults at most once, when it is first referenced. This representation is not strictly accurate, however. If a process of ten pages actually uses only half of them, then demand paging saves the I/O necessary to load the five pages that are never used. We could also increase our degree of multiprogramming by running twice as many processes. Thus,

if we had forty frames, we could run eight processes, rather than the four that could run if each required ten frames (five of which were never used).


If we increase our degree of multiprogramming, we are **over-allocating** memory. If we run six processes, each of which is ten pages in size but actually uses only five pages, we have higher CPU utilization and throughput, with ten frames to spare. It is possible, however, that each of these processes, for a particular data set, may suddenly try to use all ten of its pages, resulting in a need for sixty frames when only forty are available.

Further, consider that system memory is not used only for holding program pages. Buffers for I/O also consume a considerable amount of memory. This use can increase the strain on memory-placement algorithms. Deciding how much memory to allocate to I/O and how much to program pages is a significant challenge. Some systems allocate a fixed percentage of memory for I/O buffers, whereas others allow both processes and the I/O subsystem to compete for all system memory. Section [Missing Reference: Efficiency and performance](#) discusses the integrated relationship between I/O buffers and virtual memory techniques.

Over-allocation of memory manifests itself as follows. While a process is executing, a page fault occurs. The operating system determines where the desired page is residing on secondary storage but then finds that there are **no** free frames on the free-frame list; all memory is in use. This situation is illustrated in the animation below, where the fact that there are no free frames is depicted by a question mark.

**PARTICIPATION
ACTIVITY**

7.4.1: Need for page replacement.



Animation content:

Step 1: In this scenario, the system has two processes, eight physical memory frames, and all frames are full. Step 2: Process 1 tries to access page B, causing a page fault because the page table indicates the page is invalid. Step 3: Logical page B is on the backing store, and needs to load into physical memory for the instructions that caused the fault to complete. Step 4: Given that no physical memory frame is free, there is no frame to load B into.

Animation captions:

1. In this scenario, the system has two processes, eight physical memory frames, and all frames are full.
2. Process 1 tries to access page B, causing a page fault because the page table indicates the page is invalid.
3. Logical page B is on the backing store, and needs to load into physical memory for the instructions that caused the fault to complete.
4. Given that no physical memory frame is free, there is no frame to load B into.

The operating system has several options at this point. It could terminate the process. However, demand paging is the operating system's attempt to improve the computer system's utilization and throughput. Users should not be aware that their processes are running on a paged system—paging should be logically transparent to the user. So this option is not the best choice.

The operating system could instead use standard swapping and swap out a process, freeing all its frames and reducing the level of multiprogramming. However, as discussed in Section [Missing Reference](#) : [Swapping](#), standard swapping is no longer used by most operating systems due to the overhead of copying entire processes between memory and swap space. Most operating systems now combine swapping pages with **page replacement**, a technique we describe in detail in the remainder of this section.

Basic page replacement

Page replacement takes the following approach. If no frame is free, we find one that is not currently being used and free it. We can free a frame by writing its contents to swap space and changing the page table (and all other tables) to indicate that the page is no longer in memory (animation below). We can now use the freed frame to hold the page for which the process faulted. We modify the page-fault service routine to include page replacement:

1. Find the location of the desired page on secondary storage.
2. Find a free frame:
 - a. If there is a free frame, use it.
 - b. If there is no free frame, use a page-replacement algorithm to select a **victim frame**.

c. Write the victim frame to secondary storage (if necessary); change the page and frame tables accordingly.


3. Read the desired page into the newly freed frame; change the page and frame tables.

4. Continue the process from where the page fault occurred.

Notice that, if no frames are free, **two** page transfers (one for the page-out and one for the page-in) are required. This situation effectively doubles the page-fault service time and increases the effective access time accordingly.

PARTICIPATION ACTIVITY

7.4.2: Page Replacement.



Animation content:

Step 1: Page n is not currently in memory and has been accessed by the current instruction. A page fault occurs, and there are no free frames. Frame f is selected as the victim. Step 2: Frame f is currently holding page m. The OS writes the page to swap space on secondary storage, and frees the frame by marking m as invalid. Step 3: The OS locates page n on the backing store and schedules an I/O transfer to DMA the contents into frame f. Step 4: The transfer finishes. The OS changes the page table for n to point to frame f and that the frame is valid. The instruction is restarted and can now access the data it needs.

Animation captions:

1. Page n is not currently in memory and has been accessed by the current instruction. A page fault occurs, and there are no free frames. Frame f is selected as the victim.

2. Frame *f* is currently holding page *m*. The OS writes the page to swap space on secondary storage, and frees the frame by marking *m* as invalid.
3. The OS locates page *n* on the backing store and schedules an I/O transfer to DMA the contents into frame *f*.
4. The transfer finishes. The OS changes the page table for *n* to point to frame *f* and that the frame is valid. The instruction is restarted and can now access the data it needs.

We can reduce this overhead by using a **modify bit** (or **dirty bit**). When this scheme is used, each page or frame has a modify bit associated with it in the hardware. The modify bit for a page is set by the hardware whenever any byte in the page is written into, indicating that the page has been modified. When we select a page for replacement, we examine its modify bit. If the bit is set, we know that the page has been modified since it was read in from secondary storage. In this case, we must write the page to storage. If the modify bit is not set, however, the page has **not** been modified since it was read into memory. In this case, we need not write the memory page to storage: it is already there. This technique also applies to read-only pages (for example, pages of binary code). Such pages cannot be modified; thus, they may be discarded when desired. This scheme can significantly reduce the time required to service a page fault, since it reduces I/O time by one-half **if** the page has not been modified.

Page replacement is basic to demand paging. It completes the separation between logical memory and physical memory. With this mechanism, an enormous virtual memory can be provided for programmers on a smaller physical memory. With no demand paging, logical addresses are mapped into physical addresses, and the two sets of addresses can be different. All the pages of a process still must be in physical memory, however. With demand paging, the size of the logical address space is no longer constrained by physical memory. If we have a process of twenty pages, we can execute it in ten frames simply by using demand paging and using a replacement algorithm to find a free frame whenever necessary. If a page that has been modified is to be replaced, its contents are copied to secondary storage. A later reference to that page will cause a page fault. At that time, the page will be brought back into memory, perhaps replacing some other page in the process.

We must solve two major problems to implement demand paging: we must develop a **frame-allocation algorithm** and a **page-replacement algorithm**. That is, if we have multiple processes in memory, we must decide how many frames to allocate to each process; and when page replacement is required, we must select the frames that are to be replaced. Designing appropriate algorithms to solve these problems is an important task, because secondary storage I/O is so expensive. Even slight improvements in demand-paging methods yield large gains in system performance.

There are many different page-replacement algorithms. Every operating system probably has its own replacement scheme. How do we select a particular replacement algorithm? In general, we want the one with the lowest page-fault rate.

We evaluate an algorithm by running it on a particular string of memory references and computing the number of page faults. The string of memory references is called a **reference string**. We can generate reference strings artificially (by using a random-number generator, for example), or we can trace a given system and record the address of each memory reference. The latter choice produces a large number of data (on the order of 1 million addresses per second). To reduce the number of data, we use two facts.

First, for a given page size (and the page size is generally fixed by the hardware or system), we need to consider only the page number, rather than the entire address. Second, if we have a reference to a page p , then any references to page p that **immediately** follow will never cause a page fault. Page p will be in memory after the first reference, so the immediately following references will not fault.

For example, if we trace a particular process, we might record the following address sequence:


```
0100, 0432, 0101, 0612, 0102, 0103, 0104, 0101, 0611, 0102, 0103,  
0104, 0101, 0610, 0102, 0103, 0104, 0101, 0609, 0102, 0105
```

At 100 bytes per page, this sequence is reduced to the following reference string:

```
1, 4, 1, 6, 1, 6, 1, 6, 1, 6, 1
```

To determine the number of page faults for a particular reference string and page-replacement algorithm, we also need to know the number of page frames available. Obviously, as the number of frames available increases, the number of page faults decreases. For the reference string considered previously, for example, if we had three or more frames, we would have only three faults—one fault for the first reference to each page. In contrast, with only one frame available, we would have a replacement with every reference, resulting in eleven faults. In general, we expect a curve such as that in Figure 7.4.1. As the number of frames increases, the number of page faults drops to some minimal level. Of course, adding physical memory increases the number of frames.

Figure 7.4.1: Graph of page faults versus number of frames.



We next illustrate several page-replacement algorithms. In doing so, we use the reference string

7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

for a memory with three frames.

FIFO page replacement


The simplest page-replacement algorithm is a first-in, first-out (FIFO) algorithm. A FIFO replacement algorithm associates with each page the time when that page was brought into memory. When a page must be replaced, the oldest page is chosen. Notice that it is not strictly necessary to record the time when a page is brought in. We can create a FIFO queue to hold all pages in memory. We replace the page at the head of the queue. When a page is brought into memory, we insert it at the tail of the queue.

For our example reference string, our three frames are initially empty. The first three references (7, 0, 1) cause page faults and are brought into these empty frames. The next reference (2) replaces page 7, because page 7 was brought in first. Since 0 is the next reference and 0 is already in memory, we have no fault for this reference. The first reference to 3 results in replacement of page 0, since it is now first in line. Because of this replacement, the next reference, to 0, will fault. Page 1 is then replaced by page 0. This process continues as shown in Figure 7.4.2. Every time a fault occurs, we show which pages are in our three frames. There are fifteen faults altogether.

Figure 7.4.2: FIFO page-replacement algorithm.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

PARTICIPATION ACTIVITY

7.4.3: Mid-section review questions.



Using FIFO page replacement with demand paging, how many page faults occur with the reference string 1 2 3 4 2 3 4 1 2 1 1 3 1 4 with three page frames?

- ☐ 14
- ☐ 8
- ☐ 10

The FIFO page-replacement algorithm is easy to understand and program. However, its performance is not always good. On the one hand, the page replaced may be an initialization module that was used a long time ago and is no longer needed. On the other hand, it could contain a heavily used variable that was initialized early and is in constant use.


Notice that, even if we select for replacement a page that is in active use, everything still works correctly. After we replace an active page with a new one, a fault occurs almost immediately to retrieve the active page. Some other page must be replaced to bring the active page back into memory. Thus, a bad replacement choice increases the page-fault rate and slows process execution. It does not, however, cause incorrect execution.

To illustrate the problems that are possible with a FIFO page-replacement algorithm, consider the following reference string:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Figure 7.4.3 shows the curve of page faults for this reference string versus the number of available frames. Notice that the number of faults for four frames (ten) is **greater** than the number of faults for three frames (nine)! This most unexpected result is known as **Belady's anomaly**: for some page-replacement algorithms, the page-fault rate may **increase** as the number of allocated frames increases. We would expect that giving more memory to a process would improve its performance. In some early research, investigators noticed that this assumption was not always true. Belady's anomaly was discovered as a result.

Figure 7.4.3: Page-fault curve for FIFO replacement on a reference string.



Optimal page replacement


One result of the discovery of Belady's anomaly was the search for an **optimal page-replacement algorithm**—the algorithm that has the lowest page-fault rate of all algorithms and will never suffer from Belady's anomaly. Such an algorithm does exist and has been called OPT or MIN. It is simply this:

Replace the page that will not be used for the longest period of time.

Use of this page-replacement algorithm guarantees the lowest possible page-fault rate for a fixed number of frames.

For example, on our sample reference string, the optimal page-replacement algorithm would yield nine page faults, as shown in Figure 7.4.4. The first three references cause faults that fill the three empty frames. The reference to page 2 replaces page 7, because page 7 will not be used until reference 18, whereas page 0 will be used at 5, and page 1 at 14. The reference to page 3 replaces page 1, as page 1 will be the last of the three pages in memory to be referenced again. With only nine page faults, optimal replacement is much better than a FIFO algorithm, which results in fifteen faults. (If we ignore the first three, which all algorithms must suffer, then optimal replacement is twice as good as FIFO replacement.) In fact, no replacement algorithm can process this reference string in three frames with fewer than nine faults.

Figure 7.4.4: Optimal page-replacement algorithm.



Unfortunately, the optimal page-replacement algorithm is difficult to implement, because it requires future knowledge of the reference string. (We encountered a similar situation with the SJF CPU-scheduling algorithm in Section Shortest-job-first scheduling.) As a result, the optimal algorithm is used mainly for comparison studies. For instance, it may be useful to know that, although a new algorithm is not optimal, it is within 12.3 percent of optimal at worst and within 4.7 percent on average.

LRU page replacement

If the optimal algorithm is not feasible, perhaps an approximation of the optimal algorithm is possible. The key distinction between the FIFO and OPT algorithms (other than looking backward versus forward in time) is that the FIFO algorithm uses the time when a page was brought into memory, whereas the OPT algorithm uses the time when a page is to be **used**. If we use the recent past as an approximation of the near future, then we can replace the page that **has not been used** for the longest period of time. This approach is the **least recently used (LRU) algorithm**.

LRU replacement associates with each page the time of that page's last use. When a page must be replaced, LRU chooses the page that has not been used for the longest period of time. We can think of this strategy as the optimal page-replacement algorithm looking backward in time, rather than forward. (Strangely, if we let S^R be the reverse of a reference string S , then the page-fault rate for the OPT algorithm on S is the same as the page-fault rate for the OPT algorithm on S^R . Similarly, the page-fault rate for the LRU algorithm on S is the same as the page-fault rate for the LRU algorithm on S^R .)

The result of applying LRU replacement to our example reference string is shown in Figure 7.4.5. The LRU algorithm produces twelve faults. Notice that the first five faults are the same as those for optimal replacement. When the reference to page 4 occurs, however, LRU replacement sees that, of the three frames in memory, page 2 was used least recently. Thus, the LRU algorithm replaces page 2, not knowing that page 2 is about to be used. When it then faults for page 2, the LRU algorithm replaces page 3, since it is now the least recently used of the three pages in memory. Despite these problems, LRU replacement with twelve faults is much better than FIFO replacement with fifteen.

Figure 7.4.5: LRU page-replacement algorithm.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

PARTICIPATION ACTIVITY

7.4.4: Mid-section review questions.

Using LRU page replacement with demand paging, what is the number of page faults for the reference string 6 5 3 2 1 3 3 1 5 1 3 4 7 1 with three page frames?


- ☐ 8
- ☐ 14
- ☐ 9

The LRU policy is often used as a page-replacement algorithm and is considered to be good. The major problem is **how** to implement LRU replacement. An LRU page-replacement algorithm may require substantial hardware assistance. The problem is to determine an order for the frames defined by the time of last use. Two implementations are feasible:

- **Counters.** In the simplest case, we associate with each page-table entry a time-of-use field and add to the CPU a logical clock or counter. The clock is incremented for every memory reference. Whenever a reference to a page is made, the contents of the clock register are copied to the time-of-use field in the page-table entry for that page. In this way, we always have the "time" of the last reference to each page. We replace the page with the smallest time value. This scheme requires a search of the page table to find the LRU page and a write to memory (to the time-of-use field in the page table) for each memory access. The times must also be maintained when page tables are changed (due to CPU scheduling). Overflow of the clock must be considered.
- **Stack.** Another approach to implementing LRU replacement is to keep a stack of page numbers. Whenever a page is referenced, it is removed from the stack and put on the top. In this way, the most recently used page is always at the top of the stack, and the least recently used page is always at the bottom (Figure 7.4.6). Because entries must be removed from the middle of the stack, it is best to implement this approach by using a doubly linked list with a head pointer and a tail pointer. Removing a page and putting it on the top of the stack then requires changing six pointers at worst. Each update is a little more expensive, but there is no search for a replacement; the tail pointer

points to the bottom of the stack, which is the LRU page. This approach is particularly appropriate for software or microcode implementations of LRU replacement.

Figure 7.4.6: Use of a stack to record the most recent page references.



Like optimal replacement, LRU replacement does not suffer from Belady's anomaly. Both belong to a class of page-replacement algorithms, called **stack algorithms**, that can never exhibit Belady's anomaly. A stack algorithm is an algorithm for which it can be shown that the set of pages in memory for n frames is always a **subset** of the set of pages that would be in memory with $n + 1$ frames. For LRU replacement, the set of pages in memory would be the n most recently referenced pages. If the number of frames is increased, these n pages will still be the most recently referenced and so will still be in memory.

Note that neither implementation of LRU would be conceivable without hardware assistance beyond the standard TLB registers. The updating of the clock fields or stack must be done for **every** memory reference. If we were to use an interrupt for every reference to allow software to update such data structures, it would slow every memory reference by a factor of at least ten, hence slowing every process by a factor of ten. Few systems could tolerate that level of overhead for memory management.

LRU-approximation page replacement

Not many computer systems provide sufficient hardware support for true LRU page replacement. In fact, some systems provide no hardware support, and other page-replacement algorithms (such as a FIFO algorithm) must be used. Many systems provide some help, however, in the form of a **reference bit**. The reference bit for a page is set by the hardware whenever that page is referenced (either a read or a write to any byte in the page). Reference bits are associated with each entry in the page table.

Initially, all bits are cleared (to 0) by the operating system. As a process executes, the bit associated with each page referenced is set (to 1) by the hardware. After some time, we can determine which pages have been used and which have not been used by examining the reference bits, although we do not know the

order of use. This information is the basis for many page-replacement algorithms that approximate LRU replacement.

Additional-reference-bits algorithm

We can gain additional ordering information by recording the reference bits at regular intervals. We can keep an 8-bit byte for each page in a table in memory. At regular intervals (say, every 100 milliseconds), a timer interrupt transfers control to the operating system. The operating system shifts the reference bit for each page into the high-order bit of its 8-bit byte, shifting the other bits right by 1 bit and discarding the low-order bit. These 8-bit shift registers contain the history of page use for the last eight time periods. If the shift register contains 00000000, for example, then the page has not been used for eight time periods. A page that is used at least once in each period has a shift register value of 11111111. A page with a history register value of 11000100 has been used more recently than one with a value of 01110111. If we interpret these 8-bit bytes as unsigned integers, the page with the lowest number is the LRU page, and it can be replaced. Notice that the numbers are not guaranteed to be unique, however. We can either replace (swap out) all pages with the smallest value or use the FIFO method to choose among them.


The number of bits of history included in the shift register can be varied, of course, and is selected (depending on the hardware available) to make the updating as fast as possible. In the extreme case, the number can be reduced to zero, leaving only the reference bit itself. This algorithm is called the **second-chance page-replacement algorithm**.

Second-chance algorithm

The basic algorithm of second-chance replacement is a FIFO replacement algorithm. When a page has been selected, however, we inspect its reference bit. If the value is 0, we proceed to replace this page; but if the reference bit is set to 1, we give the page a second chance and move on to select the next FIFO page. When a page gets a second chance, its reference bit is cleared, and its arrival time is reset to the current time. Thus, a page that is given a second chance will not be replaced until all other pages have been replaced (or given second chances). In addition, if a page is used often enough to keep its reference bit set, it will never be replaced.

One way to implement the second-chance algorithm (sometimes referred to as the **clock** algorithm) is as a circular queue. A pointer (that is, a hand on the clock) indicates which page is to be replaced next. When a frame is needed, the pointer advances until it finds a page with a 0 reference bit. As it advances, it clears the reference bits (Figure 7.4.7). Once a victim page is found, the page is replaced, and the new page is inserted in the circular queue in that position. Notice that, in the worst case, when all bits are set, the pointer cycles through the whole queue, giving each page a second chance. It clears all the reference bits before selecting the next page for replacement. Second-chance replacement degenerates to FIFO replacement if all bits are set.

Figure 7.4.7: Second-chance (clock) page-replacement algorithm.



Enhanced second-chance algorithm

We can enhance the second-chance algorithm by considering the reference bit and the modify bit (described in Section Basic page replacement) as an ordered pair. With these two bits, we have the following four possible classes:

1. (0, 0) neither recently used nor modified—best page to replace
2. (0, 1) not recently used but modified—not quite as good, because the page will need to be written out before replacement
3. (1, 0) recently used but clean—probably will be used again soon
4. (1, 1) recently used and modified—probably will be used again soon, and the page will be need to be written out to secondary storage before it can be replaced

Each page is in one of these four classes. When page replacement is called for, we use the same scheme as in the clock algorithm; but instead of examining whether the page to which we are pointing has the reference bit set to 1, we examine the class to which that page belongs. We replace the first page encountered in the lowest nonempty class. Notice that we may have to scan the circular queue several times before we find a page to be replaced. The major difference between this algorithm and the simpler

clock algorithm is that here we give preference to those pages that have been modified in order to reduce the number of I/Os required.

Counting-based page replacement

There are many other algorithms that can be used for page replacement. For example, we can keep a counter of the number of references that have been made to each page and develop the following two schemes.

- The ***least frequently used*** (LFU) page-replacement algorithm requires that the page with the smallest count be replaced. The reason for this selection is that an actively used page should have a large reference count. A problem arises, however, when a page is used heavily during the initial phase of a process but then is never used again. Since it was used heavily, it has a large count and remains in memory even though it is no longer needed. One solution is to shift the counts right by 1 bit at regular intervals, forming an exponentially decaying average usage count.
- The ***most frequently used*** (MFU) page-replacement algorithm is based on the argument that the page with the smallest count was probably just brought in and has yet to be used.

As you might expect, neither MFU nor LFU replacement is common. The implementation of these algorithms is expensive, and they do not approximate OPT replacement well.

Page-buffering algorithms

Other procedures are often used in addition to a specific page-replacement algorithm. For example, systems commonly keep a pool of free frames. When a page fault occurs, a victim frame is chosen as before. However, the desired page is read into a free frame from the pool before the victim is written out. This procedure allows the process to restart as soon as possible, without waiting for the victim page to be written out. When the victim is later written out, its frame is added to the free-frame pool.

An expansion of this idea is to maintain a list of modified pages. Whenever the paging device is idle, a modified page is selected and is written to secondary storage. Its modify bit is then reset. This scheme increases the probability that a page will be clean when it is selected for replacement and will not need to be written out.

Another modification is to keep a pool of free frames but to remember which page was in each frame. Since the frame contents are not modified when a frame is written to secondary storage, the old page can be reused directly from the free-frame pool if it is needed before that frame is reused. No I/O is needed in this case. When a page fault occurs, we first check whether the desired page is in the free-frame pool. If it is not, we must select a free frame and read into it.

Some versions of the UNIX system use this method in conjunction with the second-chance algorithm. It can be a useful augmentation to any page-replacement algorithm, to reduce the penalty incurred if the wrong victim page is selected. We describe these—and other—modifications in Section Global versus local allocation.

Applications and page replacement

In certain cases, applications accessing data through the operating system's virtual memory perform worse than if the operating system provided no buffering at all. A typical example is a database, which provides its own memory management and I/O buffering. Applications like this understand their memory use and storage use better than does an operating system that is implementing algorithms for general-purpose use. Furthermore, if the operating system is buffering I/O and the application is doing so as well, then twice the memory is being used for a set of I/O.

In another example, data warehouses frequently perform massive sequential storage reads, followed by computations and writes. The LRU algorithm would be removing old pages and preserving new ones, while the application would more likely be reading older pages than newer ones (as it starts its sequential reads again). Here, MFU would actually be more efficient than LRU.

Because of such problems, some operating systems give special programs the ability to use a secondary storage partition as a large sequential array of logical blocks, without any file-system data structures. This array is sometimes called the **raw disk**, and I/O to this array is termed raw I/O. Raw I/O bypasses all the file-system services, such as file I/O demand paging, file locking, prefetching, space allocation, file names, and directories. Note that although certain applications are more efficient when implementing their own special-purpose storage services on a raw partition, most applications perform better when they use the regular file-system services.

PARTICIPATION ACTIVITY

7.4.5: Section review questions.



1) In the enhanced second chance algorithm, which of the following ordered pairs represents a page that would be the best choice for replacement?



- ☐ (0,0)
- ☐ (0,1)
- ☐ (1,0)

2) Belady's anomaly states that ____.



- ☐ for all page replacement algorithms, the page-fault rate may increase as the number of allocated frames increases
- ☐ for some page replacement algorithms, the page-fault rate may increase as the number of allocated frames increases
- ☐ for some page replacement algorithms, the page-fault rate may decrease as the number of allocated frames increases

3) Most operating systems implement _____ page replacement.



- ☐ Pure LRU
- ☐ LRU approximation

Section glossary

over-allocating: Generally, providing access to more resources than are physically available. In virtual memory, allocating more virtual memory than there is physical memory to back it.

page replacement: In virtual memory, the selection of a frame of physical memory to be replaced when a new page is allocated.

victim frame: In virtual memory, the frame selected by the page-replacement algorithm to be replaced.

modify bit: An MMU bit used to indicate that a frame has been modified (and therefore must have its contents saved before page replacement).

dirty bit: An MMU bit used to indicate that a frame has been modified (and therefore must have its contents saved before page replacement).

frame-allocation algorithm: The operating-system algorithm for allocating frames among all demands for frames.

page-replacement algorithm: In memory management, the algorithm that chooses which victim frame of physical memory will be replaced by a needed new frame of data.

reference string: A trace of accesses to a resource. In virtual memory, a list of pages accessed over a period of time.

Belady's anomaly: An anomaly in frame-allocation algorithms in which a page-fault rate may increase as the number of allocated frames increases.

optimal page-replacement algorithm: A theoretically optimal page replacement algorithm that has the lowest page-fault rate of all algorithms and never suffers from Belady's anomaly.

least recently used (LRU): In general, an algorithm that selects the item that has been used least recently. In memory management, selecting the page that has not been accessed in the longest time.

stack algorithm: A class of page-replacement algorithms that do not suffer from Belady's anomaly.

reference bit: An MMU bit indicating that a page has been referenced.

second-chance page-replacement algorithm: A FIFO page replacement algorithm in which, if the reference bit is set, the bit is cleared and the page is not replaced.

clock: In the second-chance page-replacement algorithm, a circular queue that contains possible victim frames. A frame is replaced only if it has not been recently referenced.

least frequently used (LFU): In general, an algorithm that selects the item that has been used least frequently. In virtual memory, when access counts are available, selecting the page with the lowest count.

most frequently used (MFU): In general, an algorithm that selects the item that has been used most frequently. In virtual memory, when access counts are available, selecting the page with the highest count.

raw disk: Direct access to a secondary storage device as an array of blocks with no file system.

7.5 Allocation of frames

We turn next to the issue of allocation. How do we allocate the fixed amount of free memory among the various processes? If we have 93 free frames and two processes, how many frames does each process get?

Consider a simple case of a system with 128 frames. The operating system may take 35, leaving 93 frames for the user process. Under pure demand paging, all 93 frames would initially be put on the free-frame list. When a user process started execution, it would generate a sequence of page faults. The first 93 page faults would all get free frames from the free-frame list. When the free-frame list was exhausted, a page-replacement algorithm would be used to select one of the 93 in-memory pages to be replaced with the 94th, and so on. When the process terminated, the 93 frames would once again be placed on the free-frame list.

There are many variations on this simple strategy. We can require that the operating system allocate all its buffer and table space from the free-frame list. When this space is not in use by the operating system, it can be used to support user paging. We can try to keep three free frames reserved on the free-frame list at all times. Thus, when a page fault occurs, there is a free frame available to page into. While the page swap is taking place, a replacement can be selected, which is then written to the storage device as the user process continues to execute. Other variants are also possible, but the basic strategy is clear: the user process is allocated any free frame.

Minimum number of frames

Our strategies for the allocation of frames are constrained in various ways. We cannot, for example, allocate more than the total number of available frames (unless there is page sharing). We must also allocate at least a minimum number of frames. Here, we look more closely at the latter requirement.

One reason for allocating at least a minimum number of frames involves performance. Obviously, as the number of frames allocated to each process decreases, the page-fault rate increases, slowing process execution. In addition, remember that, when a page fault occurs before an executing instruction is complete, the instruction must be restarted. Consequently, we must have enough frames to hold all the different pages that any single instruction can reference.

For example, consider a machine in which all memory-reference instructions may reference only one memory address. In this case, we need at least one frame for the instruction and one frame for the memory reference. In addition, if one-level indirect addressing is allowed (for example, a **load** instruction on frame 16 can refer to an address on frame 0, which is an indirect reference to frame 23), then paging requires at least three frames per process. (Think about what might happen if a process had only two frames.)

The minimum number of frames is defined by the computer architecture. For example, if the move instruction for a given architecture includes more than one word for some addressing modes, the instruction itself may straddle two frames. In addition, if each of its two operands may be indirect references, a total of six frames are required. As another example, the move instruction for Intel 32- and 64-bit architectures allows data to move only from register to register and between registers and memory; it does not allow direct memory-to-memory movement, thereby limiting the required minimum number of frames for a process.

Whereas the minimum number of frames per process is defined by the architecture, the maximum number is defined by the amount of available physical memory. In between, we are still left with significant choice in frame allocation.

Allocation algorithms

The easiest way to split m frames among n processes is to give everyone an equal share, m/n frames (ignoring frames needed by the operating system for the moment). For instance, if there are 93 frames and 5 processes, each process will get 18 frames. The 3 leftover frames can be used as a free-frame buffer pool. This scheme is called **equal allocation**.

An alternative is to recognize that various processes will need differing amounts of memory. Consider a system with a 1-KB frame size. If a small student process of 10 KB and an interactive database of 127 KB are the only two processes running in a system with 62 free frames, it does not make much sense to give each process 31 frames. The student process does not need more than 10 frames, so the other 21 are, strictly speaking, wasted.

To solve this problem, we can use **proportional allocation**, in which we allocate available memory to each process according to its size. Let the size of the virtual memory for process p_i be s_i , and define

$$S = \sum s_i.$$

Then, if the total number of available frames is m , we allocate a_i frames to process p_i , where a_i is approximately

$$a_i = s_i / S \times m.$$

Of course, we must adjust each a_i to be an integer that is greater than the minimum number of frames required by the instruction set, with a sum not exceeding m .

With proportional allocation, we would split 62 frames between two processes, one of 10 pages and one of 127 pages, by allocating 4 frames and 57 frames, respectively, since

$$\begin{aligned} 10/137 \times 62 &\approx 4 \text{ and} \\ 127/137 \times 62 &\approx 57. \end{aligned}$$

In this way, both processes share the available frames according to their "needs," rather than equally.

In both equal and proportional allocation, of course, the allocation may vary according to the multiprogramming level. If the multiprogramming level is increased, each process will lose some frames to provide the memory needed for the new process. Conversely, if the multiprogramming level decreases, the frames that were allocated to the departed process can be spread over the remaining processes.

Notice that, with either equal or proportional allocation, a high-priority process is treated the same as a low-priority process. By its definition, however, we may want to give the high-priority process more memory to speed its execution, to the detriment of low-priority processes. One solution is to use a proportional allocation scheme wherein the ratio of frames depends not on the relative sizes of processes but rather on the priorities of processes or on a combination of size and priority.

Global versus local allocation

Another important factor in the way frames are allocated to the various processes is page replacement. With multiple processes competing for frames, we can classify page-replacement algorithms into two broad categories: **global replacement** and **local replacement**. Global replacement allows a process to select a replacement frame from the set of all frames, even if that frame is currently allocated to some other process; that is, one process can take a frame from another. Local replacement requires that each process select from only its own set of allocated frames.

For example, consider an allocation scheme wherein we allow high-priority processes to select frames from low-priority processes for replacement. A process can select a replacement from among its own frames or the frames of any lower-priority process. This approach allows a high-priority process to increase its frame allocation at the expense of a low-priority process. Whereas with a local replacement strategy, the number of frames allocated to a process does not change, with global replacement, a process may happen to select only frames allocated to other processes, thus increasing the number of frames allocated to it (assuming that other processes do not choose **its** frames for replacement).

One problem with a global replacement algorithm is that the set of pages in memory for a process depends not only on the paging behavior of that process, but also on the paging behavior of other processes. Therefore, the same process may perform quite differently (for example, taking 0.5 seconds for one execution and 4.3 seconds for the next execution) because of totally external circumstances. Such is not the case with a local replacement algorithm. Under local replacement, the set of pages in memory for a process is affected by the paging behavior of only that process. Local replacement might hinder a process, however, by not making available to it other, less used pages of memory. Thus, global replacement generally results in greater system throughput. It is therefore the more commonly used method.

Major and minor page faults

As described in Section Basic concepts, a page fault occurs when a page does not have a valid mapping in the address space of a process. Operating systems generally distinguish between two types of page faults: **major** and **minor** faults. (Windows refers to major and minor faults as **hard** and **soft** faults, respectively.) A major page fault occurs when a page is referenced and the page is not in memory. Servicing a major page fault requires reading the desired page from the backing store into a free frame and updating the page table. Demand paging typically generates an initially high rate of major page faults.

Minor page faults occur when a process does not have a logical mapping to a page, yet that page is in memory. Minor faults can occur for one of two reasons. First, a process may reference a shared library that is in memory, but the process does not have a mapping to it in its page table. In this instance, it is only necessary to update the page table to refer to the existing page in memory. A second cause of minor faults occurs when a page is reclaimed from a process and placed on the free-frame list, but the page has not yet been zeroed out and allocated to another process. When this kind of fault occurs, the frame is removed from the free-frame list and reassigned to the process. As might be expected, resolving a minor page fault is typically much less time consuming than resolving a major page fault.

You can observe the number of major and minor page faults in a Linux system using the command `ps -eo min_flt,maj_flt,cmd`, which outputs the number of minor and major page faults, as well as the command that launched the process. An example output of this `ps` command appears below:


MINFL	MAJFL	CMD
186509	32	/usr/lib/systemd/systemd-logind
76822	9	/usr/sbin/sshd -D
1937	0	vim 10.tex
699	14	/sbin/auditd -n

Here, it is interesting to note that, for most commands, the number of major page faults is generally quite low, whereas the number of minor faults is much higher. This indicates that Linux processes likely take significant advantage of shared libraries as, once a library is loaded in memory, subsequent page faults are only minor faults.

Next, we focus on one possible strategy that we can use to implement a global page-replacement policy. With this approach, we satisfy all memory requests from the free-frame list, but rather than waiting for the list to drop to zero before we begin selecting pages for replacement, we trigger page replacement when the list falls below a certain threshold. This strategy attempts to ensure there is always sufficient free memory to satisfy new requests.

Such a strategy is depicted in Figure 7.5.1. The strategy's purpose is to keep the amount of free memory above a minimum threshold. When it drops below this threshold, a kernel routine is triggered that begins reclaiming pages from all processes in the system (typically excluding the kernel). Such kernel routines are often known as **reapers**, and they may apply any of the page-replacement algorithms covered in Section 7.4. When the amount of free memory reaches the maximum threshold, the reaper routine is suspended, only to resume once free memory again falls below the minimum threshold.

Figure 7.5.1: Reclaiming pages.



In Figure [7.5.1](#), we see that at point a the amount of free memory drops below the minimum threshold, and the kernel begins reclaiming pages and adding them to the free-frame list. It continues until the maximum threshold is reached (point b). Over time, there are additional requests for memory, and at point c the amount of free memory again falls below the minimum threshold. Page reclamation resumes, only to be suspended when the amount of free memory reaches the maximum threshold (point d). This process continues as long as the system is running.

As mentioned above, the kernel reaper routine may adopt any page-replacement algorithm, but typically it uses some form of LRU approximation. Consider what may happen if the reaper routine is unable to maintain the list of free frames above the minimum threshold. Under these circumstances, the reaper routine may begin to reclaim pages more aggressively. For example, perhaps it will suspend the second-chance algorithm and use pure FIFO. Another, more extreme, example occurs in Linux; when the amount of free memory falls to **very** low levels, a routine known as the **out-of-memory (OOM) killer** selects a process to terminate, thereby freeing its memory. How does Linux determine which process to terminate? Each process has what is known as an OOM score, with a higher score increasing the likelihood that the process could be terminated by the OOM killer routine. OOM scores are calculated according to the percentage of memory a process is using—the higher the percentage, the higher the OOM score. (OOM


scores can be viewed in the `/proc` file system, where the score for a process with `pid` 2500 can be viewed as `/proc/2500/oom_score`.)

In general, not only can reaper routines vary how aggressively they reclaim memory, but the values of the minimum and maximum thresholds can be varied as well. These values can be set to default values, but some systems may allow a system administrator to configure them based on the amount of physical memory in the system.

Non-uniform memory access

Thus far in our coverage of virtual memory, we have assumed that all main memory is created equal—or at least that it is accessed equally. On **non-uniform memory access (NUMA)** systems with multiple CPUs (Section Multiprocessor systems), that is not the case. On these systems, a given CPU can access some sections of main memory faster than it can access others. These performance differences are caused by how CPUs and memory are interconnected in the system. Such a system is made up of multiple CPUs, each with its own local memory (Figure 7.5.2). The CPUs are organized using a shared system interconnect, and as you might expect, a CPU can access its local memory faster than memory local to another CPU. NUMA systems are without exception slower than systems in which all accesses to main memory are treated equally. However, as described in Section Multiprocessor systems, NUMA systems can accommodate more CPUs and therefore achieve greater levels of throughput and parallelism.

Figure 7.5.2: NUMA multiprocessing architecture.



Managing which page frames are stored at which locations can significantly affect performance in NUMA systems. If we treat memory as uniform in such a system, CPUs may wait significantly longer for memory access than if we modify memory allocation algorithms to take NUMA into account. We described some of these modifications in Section Processor affinity. Their goal is to have memory frames allocated "as

close as possible" to the CPU on which the process is running. (The definition of **close** is "with minimum latency," which typically means on the same system board as the CPU). Thus, when a process incurs a page fault, a NUMA-aware virtual memory system will allocate that process a frame as close as possible to the CPU on which the process is running.

To take NUMA into account, the scheduler must track the last CPU on which each process ran. If the scheduler tries to schedule each process onto its previous CPU, and the virtual memory system tries to allocate frames for the process close to the CPU on which it is being scheduled, then improved cache hits and decreased memory access times will result.

The picture is more complicated once threads are added. For example, a process with many running threads may end up with those threads scheduled on many different system boards. How should the memory be allocated in this case?

As we discussed in Section Example: Linux scheduling, Linux manages this situation by having the kernel identify a hierarchy of scheduling domains. The Linux CFS scheduler does not allow threads to migrate across different domains and thus incur memory access penalties. Linux also has a separate free-frame list for each NUMA node, thereby ensuring that a thread will be allocated memory from the node on which it is running. Solaris solves the problem similarly by creating **lggroups** (for "locality groups") in the kernel. Each lggroup gathers together CPUs and memory, and each CPU in that group can access any memory in the group within a defined latency interval. In addition, there is a hierarchy of lggroups based on the amount of latency between the groups, similar to the hierarchy of scheduling domains in Linux. Solaris tries to schedule all threads of a process and allocate all memory of a process within an lggroup. If that is not possible, it picks nearby lggroups for the rest of the resources needed. This practice minimizes overall memory latency and maximizes CPU cache hit rates.

**PARTICIPATION
ACTIVITY**

7.5.1: Section review questions.



1) _____ allocation reclaims pages from any process in memory, not just the process incurring a page fault.



- ☐ local
- ☐ global

2) Non-uniform memory access (NUMA)



- ☐ exists when memory access times can vary
- ☐ refers to the hierarchy of memory from registers through main memory to secondary storage.
- ☐ occurs when page references are resolved in the TLB without having to access the page table.

Section glossary

equal allocation: An allocation algorithm that assigns equal amounts of a resource to all requestors. In virtual memory, assigning an equal number of frames to each process.

proportional allocation: An allocation algorithm that assigns a resource in proportion to some aspect of the requestor. In virtual memory, the assignment of page frames in proportion to the size of each process.

global replacement: In virtual memory frame allocation, allowing a process to select a replacement frame from the set of all frames in the system, not just those allocated to the process.

local replacement: In virtual-memory frame allocation, allowing a process to select a replacement frame only from the set of all frames allocated to the process.

major fault: In virtual memory, a page fault resolved by executing I/O to bring in the page from secondary storage.

minor fault: In virtual memory, a page fault that can be resolved without having to page in data from secondary storage.

reapers: In memory management, routines that scan memory, freeing frames to maintain a minimum level of available free memory.

out-of-memory (OOM) killer: In Linux, a routine that executes when free memory is very low, terminating processes to free memory.

non-uniform memory access (NUMA): An architecture aspect of many computer systems in which the time to access memory varies based on which core the thread is running on (e.g., a core interlink is slower than accessing DIMMs directly attached to core).

lgroups: In Solaris, locality groups located in the kernel; each lgroup gathers together CPUs and memory, and each CPU in that group can access any memory in the group within a defined latency interval. A method for dealing with NUMA.

7.6 Thrashing

Consider what occurs if a process does not have "enough" frames—that is, it does not have the minimum number of frames it needs to support pages in the working set. The process will quickly page-fault. At this point, it must replace some page. However, since all its pages are in active use, it must replace a page that will be needed again right away. Consequently, it quickly faults again, and again, and again, replacing pages that it must bring back in immediately.

This high paging activity is called **thrashing**. A process is thrashing if it is spending more time paging than executing. As you might expect, thrashing results in severe performance problems.


Cause of thrashing

Consider the following scenario, which is based on the actual behavior of early paging systems. The operating system monitors CPU utilization. If CPU utilization is too low, we increase the degree of multiprogramming by introducing a new process to the system. A global page-replacement algorithm is used; it replaces pages without regard to the process to which they belong. Now suppose that a process enters a new phase in its execution and needs more frames. It starts faulting and taking frames away from other processes. These processes need those pages, however, and so they also fault, taking frames from other processes. These faulting processes must use the paging device to swap pages in and out. As they queue up for the paging device, the ready queue empties. As processes wait for the paging device, CPU utilization decreases.

The CPU scheduler sees the decreasing CPU utilization and **increases** the degree of multiprogramming as a result. The new process tries to get started by taking frames from running processes, causing more page faults and a longer queue for the paging device. As a result, CPU utilization drops even further, and the CPU scheduler tries to increase the degree of multiprogramming even more. Thrashing has occurred, and system throughput plunges. The page-fault rate increases tremendously. As a result, the effective memory-access time increases. No work is getting done, because the processes are spending all their time paging.

This phenomenon is illustrated in Figure [7.6.1](#), in which CPU utilization is plotted against the degree of multiprogramming. As the degree of multiprogramming increases, CPU utilization also increases, although more slowly, until a maximum is reached. If the degree of multiprogramming is increased further, thrashing sets in, and CPU utilization drops sharply. At this point, to increase CPU utilization and stop thrashing, we must **decrease** the degree of multiprogramming.

Figure 7.6.1: Thrashing.




We can limit the effects of thrashing by using a **local replacement algorithm** (or **priority replacement algorithm**). As mentioned earlier, local replacement requires that each process select from only its own set of allocated frames. Thus, if one process starts thrashing, it cannot steal frames from another process and cause the latter to thrash as well. However, the problem is not entirely solved. If processes are thrashing, they will be in the queue for the paging device most of the time. The average service time for a page fault will increase because of the longer average queue for the paging device. Thus, the effective access time will increase even for a process that is not thrashing.

To prevent thrashing, we must provide a process with as many frames as it needs. But how do we know how many frames it "needs"? One strategy starts by looking at how many frames a process is actually using. This approach defines the **locality model** of process execution.

The locality model states that, as a process executes, it moves from locality to locality. A locality is a set of pages that are actively used together. A running program is generally composed of several different localities, which may overlap. For example, when a function is called, it defines a new locality. In this locality, memory references are made to the instructions of the function call, its local variables, and a subset of the global variables. When we exit the function, the process leaves this locality, since the local variables and instructions of the function are no longer in active use. We may return to this locality later.

Figure 7.6.2 illustrates the concept of locality and how a process's locality changes over time. At time (a), the locality is the set of pages {18, 19, 20, 21, 22, 23, 24, 29, 30, 33}. At time (b), the locality changes to {18, 19, 20, 24, 25, 26, 27, 28, 29, 31, 32, 33}. Notice the overlap, as some pages (for example, 18, 19, and 20) are part of both localities.

Figure 7.6.2: Locality in a memory-reference pattern.




Thus, we see that localities are defined by the program structure and its data structures. The locality model states that all programs will exhibit this basic memory reference structure. Note that the locality model is the unstated principle behind the caching discussions so far in this book. If accesses to any types of data were random rather than patterned, caching would be useless.

Suppose we allocate enough frames to a process to accommodate its current locality. It will fault for the pages in its locality until all these pages are in memory; then, it will not fault again until it changes localities. If we do not allocate enough frames to accommodate the size of the current locality, the process will thrash, since it cannot keep in memory all the pages that it is actively using.

Working-set model

The **working-set model** is based on the assumption of locality. This model uses a parameter, Δ , to define the **working-set window**. The idea is to examine the most recent Δ page references. The set of pages in the most recent Δ page references is the **working set** (Figure 7.6.3). If a page is in active use, it will be in the working set. If it is no longer being used, it will drop from the working set Δ time units after its last reference. Thus, the working set is an approximation of the program's locality.

Figure 7.6.3: Working-set model.



For example, given the sequence of memory references shown in Figure 7.6.3, if $\Delta = 10$ memory references, then the working set at time t_1 is $\{1, 2, 5, 6, 7\}$. By time t_2 , the working set has changed to $\{3, 4\}$.

The accuracy of the working set depends on the selection of Δ . If Δ is too small, it will not encompass the entire locality; if Δ is too large, it may overlap several localities. In the extreme, if Δ is infinite, the working set is the set of pages touched during the process execution.

The most important property of the working set, then, is its size. If we compute the working-set size, WSS_i , for each process in the system, we can then consider that

$$D = \sum WSS_i,$$

where D is the total demand for frames. Each process is actively using the pages in its working set. Thus, process i needs WSS_i frames. If the total demand is greater than the total number of available frames ($D > m$), thrashing will occur, because some processes will not have enough frames.

Once Δ has been selected, use of the working-set model is simple. The operating system monitors the working set of each process and allocates to that working set enough frames to provide it with its working-set size. If there are enough extra frames, another process can be initiated. If the sum of the working-set sizes increases, exceeding the total number of available frames, the operating system selects a process to suspend. The process's pages are written out (swapped), and its frames are reallocated to other processes. The suspended process can be restarted later.

This working-set strategy prevents thrashing while keeping the degree of multiprogramming as high as possible. Thus, it optimizes CPU utilization. The difficulty with the working-set model is keeping track of the working set. The working-set window is a moving window. At each memory reference, a new reference appears at one end, and the oldest reference drops off the other end. A page is in the working set if it is referenced anywhere in the working-set window.

**PARTICIPATION
ACTIVITY**

7.6.1: Mid-section review question.

OpenStax, 02/10/2018 05:21:56 AM
Admin Tool



_____ occurs when a process spends more time paging than executing.




Check

[Show answer](#)

Working sets and page-fault rates

There is a direct relationship between the working set of a process and its page-fault rate. Typically, as shown in Figure 7.6.3, the working set of a process changes over time as references to data and code sections move from one locality to another. Assuming there is sufficient memory to store the working set of a process (that is, the process is not thrashing), the page-fault rate of the process will transition between peaks and valleys over time. This general behavior is shown below:



A peak in the page-fault rate occurs when we begin demand-paging a new locality. However, once the working set of this new locality is in memory, the page-fault rate falls. When the process

moves to a new working set, the page-fault rate rises toward a peak once again, returning to a lower rate once the new working set is loaded into memory. The span of time between the start of one peak and the start of the next peak represents the transition from one working set to another.

We can approximate the working-set model with a fixed-interval timer interrupt and a reference bit. For example, assume that Δ equals 10,000 references and that we can cause a timer interrupt every 5,000 references. When we get a timer interrupt, we copy and clear the reference-bit values for each page. Thus, if a page fault occurs, we can examine the current reference bit and two in-memory bits to determine whether a page was used within the last 10,000 to 15,000 references. If it was used, at least one of these bits will be on. If it has not been used, these bits will be off. Pages with at least one bit on will be considered to be in the working set.


Note that this arrangement is not entirely accurate, because we cannot tell where, within an interval of 5,000, a reference occurred. We can reduce the uncertainty by increasing the number of history bits and the frequency of interrupts (for example, 10 bits and interrupts every 1,000 references). However, the cost to service these more frequent interrupts will be correspondingly higher.

Page-fault frequency

The working-set model is successful, and knowledge of the working set can be useful for prepaging (Section Prepaging), but it seems a clumsy way to control thrashing. A strategy that uses the **page-fault frequency (PFF)** takes a more direct approach.

The specific problem is how to prevent thrashing. Thrashing has a high page-fault rate. Thus, we want to control the page-fault rate. When it is too high, we know that the process needs more frames. Conversely, if the page-fault rate is too low, then the process may have too many frames. We can establish upper and lower bounds on the desired page-fault rate (Figure 7.6.4). If the actual page-fault rate exceeds the upper limit, we allocate the process another frame. If the page-fault rate falls below the lower limit, we remove a frame from the process. Thus, we can directly measure and control the page-fault rate to prevent thrashing.

Figure 7.6.4: Page-fault frequency.



As with the working-set strategy, we may have to swap out a process. If the page-fault rate increases and no free frames are available, we must select some process and swap it out to backing store. The freed frames are then distributed to processes with high page-fault rates.

Current practice

Practically speaking, thrashing and the resulting swapping have a disagreeably high impact on performance. The current best practice in implementing a computer system is to include enough physical memory, whenever possible, to avoid thrashing and swapping. From smartphones through large servers, providing enough memory to keep all working sets in memory concurrently, except under extreme conditions, provides the best user experience.

PARTICIPATION ACTIVITY

7.6.2: Section review questions.

1) The _____ refers to the set of pages a process references over a period of time.

- ☐ reference string
- ☐ working set
- ☐ page fault rate

2) What best describes the relationship between the working set and page fault rate of a process?

- ☐ The page fault rate is independent of the working set.
- ☐ Once the working set of a process is in memory, the page fault rate will be low.
- ☐ The larger the working set the higher the page fault rate.

Section glossary

thrashing: Paging memory at a high rate. A system thrashes when there is insufficient physical memory to meet virtual memory demand.

local replacement algorithm: A virtual-memory page replacement algorithm that avoids thrashing by not allowing a process to steal frames from other processes.

priority replacement algorithm: A virtual memory page replacement algorithm that avoids thrashing by not allowing a process to steal frames from other processes.

locality model: A model for page replacement based on the working-set strategy.

working-set model: A model of memory access based on tracking the set of most recently accessed pages.

working-set window: A limited set of most recently accessed pages (a "window" view of the entire set of accessed pages).

working set: The set of pages in the most recent page references.

page-fault frequency: The frequency of page faults.


7.7 Memory compression

An alternative to paging is **memory compression**. Here, rather than paging out modified frames to swap space, we compress several frames into a single frame, enabling the system to reduce memory usage without resorting to swapping pages.

In Figure [7.7.1](#), the free-frame list contains six frames. Assume that this number of free frames falls below a certain threshold that triggers page replacement. The replacement algorithm (say, an LRU approximation


algorithm) selects four frames—15, 3, 35, and 26—to place on the free-frame list. It first places these frames on a modified-frame list. Typically, the modified-frame list would next be written to swap space, making the frames available to the free-frame list. An alternative strategy is to compress a number of frames—say, three—and store their compressed versions in a single page frame.

Figure 7.7.1: Free-frame list before compression.



In Figure 7.7.2, frame 7 is removed from the free-frame list. Frames 15, 3, and 35 are compressed and stored in frame 7, which is then stored in the list of compressed frames. The frames 15, 3, and 35 can now be moved to the free-frame list. If one of the three compressed frames is later referenced, a page fault occurs, and the compressed frame is decompressed, restoring the three pages 15, 3, and 35 in memory.

Figure 7.7.2: Free-frame list after compression



As we have noted, mobile systems generally do not support either standard swapping or swapping pages. Thus, memory compression is an integral part of the memory-management strategy for most mobile operating systems, including Android and iOS. In addition, both Windows 10 and macOS support memory compression. For Windows 10, Microsoft developed the **Universal Windows Platform (UWP)** architecture, which provides a common app platform for devices that run Windows 10, including mobile devices. UWP apps running on mobile devices are candidates for memory compression. macOS first supported memory compression with Version 10.9 of the operating system, first compressing LRU pages when free memory is short and then paging if that doesn't solve the problem. Performance tests indicate that memory compression is faster than paging even to SSD secondary storage on laptop and desktop macOS systems.

Although memory compression does require allocating free frames to hold the compressed pages, a significant memory saving can be realized, depending on the reductions achieved by the compression algorithm. (In the example above, the three frames were reduced to one-third of their original size.) As with any form of data compression, there is contention between the speed of the compression algorithm and the amount of reduction that can be achieved (known as the **compression ratio**). In general, higher compression ratios (greater reductions) can be achieved by slower, more computationally expensive algorithms. Most algorithms in use today balance these two factors, achieving relatively high compression ratios using fast algorithms. In addition, compression algorithms have improved by taking advantage of multiple computing cores and performing compression in parallel. For example, Microsoft's Xpress and Apple's WKdm compression algorithms are considered fast, and they report compressing pages to 30 to 50 percent of their original size.

**PARTICIPATION
ACTIVITY**

7.7.1: Section review questions.



What must occur if a compressed frame is referenced?



- ☐ The contents of the compressed frame must be accessed.
- ☐ The frame must be decompressed and moved to a free frame.

Section glossary

memory compression: In memory management, an alternative to paging involving compressing the contents of frames to decrease the memory used.

universal Windows platform (UWP): Windows 10 architecture that provides a common app platform for all devices that run it, including mobile devices.

compression ratio: In memory compression, a measurement of the effectiveness of the compression (the ratio of the compressed space to the original amount of uncompressed space).

7.8 Principles of virtual memory

Demand paging

Virtual memory (VM) is a collection of one or more logical address spaces, each of which may exceed the size of physical memory. A logical address is then referred to as a virtual address. A paged VM creates a single large contiguous address space per process. A paged VM with segmentation creates multiple large address spaces per process, each of which is paged.

Demand paging is the principle of loading a page into memory only when the page is needed, rather than at the start of the execution.


A **present bit** is a binary flag in each page table entry that indicates whether the corresponding page is currently resident in memory. If a page is resident, then the entry points to the frame that holds the page.

A **page fault** is an interrupt that occurs when a program attempts to reference a non-resident page. The interrupt triggers the operating system to find the page on disk, copy the page into a frame, and set the present bit to 1.

Demand paging makes the implementation of VM possible by moving pages to main memory only as needed.

PARTICIPATION ACTIVITY

7.8.1: Demand paging.



Animation content:

Static figure: A conceptual diagram of virtual memory (VM) management using demand paging. The figure includes three main components: a virtual memory (VM) column, a page table, and a physical memory column. The VM column shows a series of memory pages numbered from 0 to $n-1$, with a partition indicating the occupied space kept on disk, which ranges from page 0 to page r . The page table is a representation of the mapping of the VM pages to the physical memory frames. Each entry in the page table shows whether the corresponding VM page is present in physical memory, indicated by a 1, or not, indicated by a 0, with an arrow pointing to the respective frame in physical memory if the page is resident. A label present bit points to the bits that indicate the presence of a page. The physical memory column illustrates the frames of memory filled with the contents of the corresponding VM pages.

Step 1: A VM is a large address space consisting of pages $[0:n-1]$. The portion occupied by the program (pages $[0:r]$) is kept on disk. The remaining pages are just an abstraction.

Step 2: A page table keeps track of which pages are currently in physical memory. If a page is resident, then the present bit is set and the entry field points to the corresponding frame.

Pages 2 and k are currently resident in frames 2 and 5, respectively. The two frames are highlighted.

Step 3: If execution crosses a page boundary into a non-resident page, the present bit triggers a page fault, which causes the page to be copied into a free frame. The page table is updated.

Frame 0 is highlighted, the present bit in page table entry 3 is set to 1, and the entry points to frame 0.

Step 4: Similarly, when execution references a non-resident data page, a page fault causes the page to be copied into memory and the page table to be updated accordingly. A marker points to page r in VM. Frame 4 is highlighted, the present bit in page table entry r is set to 1, and the entry points to frame r .

Animation captions:

1. A VM is a large address space consisting of pages $[0:n-1]$. The portion occupied by the program (pages $[0:r]$) is kept on disk. The remaining pages are just an abstraction.
2. A page table keeps track of which pages are currently in physical memory. If a page is resident, then the present bit is set and the entry field points to the corresponding frame.
3. If execution crosses a page boundary into a non-resident page, the present bit triggers a page fault, which causes the page to be copied into a free frame. The page table is updated.
4. Similarly, when execution references a non-resident data page, a page fault causes the page to be copied into memory and the page table to be updated accordingly.



1) When $p = 0$ then the corresponding page resides ____.



- ☐ in physical memory only
- ☐ on disk only
- ☐ in both physical memory and on disk

2) When $p = 1$ then the corresponding page resides ____.



- ☐ in physical memory only
- ☐ on disk only
- ☐ in both physical memory and on disk



A process executes the code

```
br m
```

```
...
```

```
m: ld R, x
```

where the branch instruction is on a resident page and the instruction "ld" loads the content of the variable x into a register R.

Executing the code ____.



- ☐ cannot trigger any page faults
- ☐ can trigger at most 1 page fault
- ☐ can trigger 0, 1, or 2 page faults

Page replacement

Virtual memory is typically much larger than the available physical memory. When all frames are occupied and a page fault occurs, one of the resident pages must be removed from memory to create a free frame.

Page replacement is the act of overwriting a page in memory with a different page loaded from the disk when needed.


Moving pages between memory and the disk is very time consuming and must be minimized. When the page to be removed has not been modified during execution then an exact copy still exists on the disk and the page does not have to be copied back to the disk.

A **modified-bit** (**m-bit**) is a binary flag in each page table entry that indicates whether the corresponding page has been modified during execution. The modified bit is set to 1 automatically by any instruction that stores data into the page and is used by the operating system to minimize the movement of data to disk.

PARTICIPATION ACTIVITY

7.8.4: Page replacement.

© 2011 Pearson Education, Inc.
All rights reserved.



Animation content:

Static figure: A conceptual illustration of virtual memory, a page table, and physical memory. Virtual memory (VM) is depicted as a column with cells labeled from 0 to $n-1$, indicating the range of virtual page numbers. Each cell represents a page, and the occupied space is marked as kept on disk. The page table is a representation of the mapping of the VM pages to the physical memory frames. Each entry in the page table has 3 components: a frame number, a present bit, and a modified bit. The present bit indicates whether the page is in physical memory, while the modified bit indicates whether the page has been altered during execution. If the present bit is 1, then the frame number indicates in which physical memory frame the page is kept. Physical memory is depicted as a column with cells

labeled 0 through 5. The frames occupied by virtual pages are highlighted. Arrows point from occupied page table entries the corresponding frames in physical memory.

Step 1: The page table uses the present bit to keep track of which pages are in physical memory, and the modified bit to keep track of which pages have been modified. Page table entries 2, 3, and k point to frames 2, 0, and 5, respectively. The corresponding present bits are 1 and the modified bits are 0.

Step 2: As the program executes, some of the memory-resident pages become modified, making the copy on disk obsolete (page k). The modified bit is set to 1.

Step 3: When a page fault decides to replace page k in frame 5 with page r then page k must first be written back to disk before being overwritten by page r. Page k in frame 5 is copied back to page k in VM. Then page r is copied from VM to frame 5 and the page table is updated. Entry r points to frame 5, the present bit is 1 and the modified bit is 0.

Step 4: When a page fault decides to replace page 2 with page 0 in frame 2 then page 2 can be overwritten without saving because an exact copy is still on disk. Page 0 is copied from VM to frame 2 and the page table is updated. Entry 2 points to frame 0, the present bit is 1 and the modified bit is 0.

Animation captions:

1. The page table uses the present bit to keep track of which pages are in physical memory, and the modified bit to keep track of which pages have been modified.
2. As the program executes, some of the memory-resident pages become modified, making the copy on disk obsolete (page k). The modified bit is set to 1.
3. When a page fault decides to replace page k in frame 5 with page r then page k must first be written back to disk before being overwritten by page r.
4. When a page fault decides to replace page 2 with page 0 in frame 2 then page 2 can be overwritten without saving because an exact copy is still on disk.

PARTICIPATION ACTIVITY

7.8.5: Requirements for VM implementation.



1) VM cannot be implemented without demand paging.



- ☐ True
☐ False

2) VM cannot be implemented without page replacement.



- ☐ True
☐ False

3) VM cannot be implemented without the modified bit.

- ☐ True
- ☐ False

**PARTICIPATION
ACTIVITY**

7.8.6: The significance of the present bit (p) and the modified bit (m).

p = 0 means the page is currently not resident in memory.

p = 1 means the page is currently resident in memory.

m = 0 means the page has not been modified.

m = 1 means the page has been modified.

1) A page must be written back to disk when ____.

- ☐ p = 0 and m = 0
- ☐ p = 1 and m = 0
- ☐ p = 1 and m = 1

2) Immediately after a page replacement, the bits in the page table are ____.

- ☐ p = 0 and m = 0
- ☐ p = 1 and m = 0
- ☐ p = 1 and m = 1

3) A page fault is triggered when ____.

- ☐ p = 0 and m = 0
- ☐ p = 1 and m = 0
- ☐ p = 1 and m = 1

Paging of tables


Virtual memory can create address spaces much larger than the available physical memory. Ex: With a 64-bit virtual address, the potential size of VM is $2^{64} = 18,446,744,073,709,551,616$ words. With 12 bits for the offset w , the remaining 52 bits, when split evenly between s and p , would result in segment and page tables of $2^{26} = 67,108,864$ entries. An uneven split would make one of the tables even larger.

Since few programs make use of such large address spaces, much of the space taken up by a table in memory would be wasted. To reduce the overhead, a table can be divided into individual pages and brought into memory as needed using demand paging and page replacement.

A segment table can be divided into pages by dividing the segment number s into two components, s_1 and s_2 , and to use a new page table to keep track of the pages constituting the segment table. Similarly, a large page table can be divided into pages and kept track of by another higher-order page table.

PARTICIPATION ACTIVITY

7.8.7: Address translation with a paged page table.



Animation content:

Static figure: A 4-level horizontal hierarchy shows the relationship between (1) a segment table, (2) a page table PT' of page table PT , (3) a page table PT of segment s , and (4) the pages of segment s in physical memory. Vertical arrows show offsets within each table and dashed lines connect table entries from one column to tables or pages in the next column.

A virtual address translation $(s, p, w) \Rightarrow (s, p_1, p_2, w)$ is shown below.

Step 1: A 3 level hierarchy starts with a segment table. The entry s points to the page table of segment s and each entry of the page table points to one page of segment s in physical memory. Pages 1, 2,

and k of segment s are shown in physical memory.

Step 2: A page table that is too large to keep in memory may be subdivided into pages, each entry still pointing to the same pages as before. The page table is subdivided into pages 0, 1, ..., p_1 .

Step 3: To keep track of the pages constituting the page table, a higher order page table PT' of the page table PT is introduced. Each entry points to a page of the page table. A new table, Page table PT' of page table PT , is interposed between the segment table and the individual pages of PT .

Step 4: The page number p is divided into two components, p_1 and p_2 , resulting in a virtual address of the form (s, p_1, p_2, w) . The address translation $(s, p, w) \Rightarrow (s, p_1, p_2, w)$ appears

Step 5: s selects the segment. p_1 selects the entry in the page table PT' of the page table PT that points to page p_1 of PT .

Step 6: p_2 is an offset into page p_1 of PT , which contains the frame number of the page containing the word corresponding to the virtual address (s, p_1, p_2, w) .

Animation captions:

1. A 3 level hierarchy starts with a segment table. The entry s points to the page table of segment s and each entry of the page table points to one page of segment s in physical memory.
2. A page table that is too large to keep in memory may be subdivided into pages, each entry still pointing to the same pages as before.
3. To keep track of the pages constituting the page table, a higher order page table PT' of the page table PT is introduced. Each entry points to a page of the page table.
4. The page number p is divided into two components, p_1 and p_2 , resulting in a virtual address of the form (s, p_1, p_2, w) .
5. s selects the segment. p_1 selects the entry in the page table PT' of the page table PT that points to page p_1 of PT .
6. p_2 is an offset into page p_1 of PT , which contains the frame number of the page containing the word corresponding to the virtual address (s, p_1, p_2, w) .

PARTICIPATION ACTIVITY

7.8.8: Division of VM where the page tables are paged.



The VM addresses use:

s : 5 bits

p_1 : 3 bits

p_2 : 4 bits

1) The number of segments is ____.

☐ 2^3

☐ 2^4

☐ 2^5

2) The number of pages per segment is ____.

☐ 2^4

☐ 2^5

☐ 2^7

**PARTICIPATION
ACTIVITY**

7.8.9: Division of VM where the segment table is paged.

The VM addresses use:

s1: 3 bits

s2: 4 bits

p: 5 bits

1) The number of segments is ____.

☐ 2^3

☐ 2^4

☐ 2^7

2) The number of pages per segment is ____.

☐ 2^3

☐ 2^4

☐ 2^5

**PARTICIPATION
ACTIVITY**

7.8.10: Division of VM where both the segment table and the page tables are paged.

The VM address uses:

s1: 3 bits

s2: 3 bits

p1: 3 bits

p2: 3 bits

w: 8 bits

The total number of words that make up the VM is ____.

☐ 2^8

☐ 2^{12}

☐ 2^{20}

**CHALLENGE
ACTIVITY**

7.8.1: Virtual address translation.

726670.4301284.qx3zqy7

Start

Assume a virtual memory that uses segmentation and paging for a single process.

Segment table and page table entries contain frame numbers of the corresponding page tables or pag

Size of segment table = size of page table = size of page = size of frame = 512 words.

Segment table (ST) starts at address 0 of physical memory.

Physical memory (PM) has the following contents:

	0	1	2	3	4	...
PM	4	7	8	5	3	...

The page table (PT) of segment 2 starts at physical address

The page table (PT) of segment 3 starts at physical address

1

2

3

4

5

Check

Next



EXERCISE

7.8.1: Demand paging.



A paged memory is shared by two processes, p1 and p2.

0	6	8	9	16	0	24	5
1	15	9	-3	17	-12	25	-3
2	3	10	8	18	-16	26	-3
3	-2	11	12	19	---	27	2
4	-45	12	0	20	0	28	0
5	8	13	-12	21	0	29	1
6	28	14	24	22	0	30	2
7	-2	15	---	23	0	31	3

Page table
of p1Page table
of p2

- The size of each page and page table is 4.
 - The page table of p1 begins at address 4. The page table of p2 begins at address 12.
 - A dash denotes a nonexistent page.
 - The sign bit serves as the present bit. Thus a positive number in a page table is the address of a resident page and a negative number denotes a non-resident page.
- (a) Both processes reference the virtual addresses: 0, 1, 4, 5, 8, 10, 12, 15.
For each process, translate the virtual addresses into the corresponding physical addresses. If the translation is successful, read the content of the corresponding memory location. Otherwise, indicate whether a page fault or an error has occurred. (Assume that page faults and errors are not resolved and thus all of the virtual addresses are independent from each other.)



EXERCISE

7.8.2: Paged segment table.



A 32-bit virtual address is split into (s, p, w), where p and w are both 9 bits and s is 14 bits. If the segment table is paged, the 14 bits of s are split into s1 and s2.

- (a) Which choice is preferable:
1. s1 is 9 bits and s2 is 5 bits
 2. s1 is 5 bits and s2 is 9 bits
- (b) How does the choice impact the total size of virtual memory?

**EXERCISE**

7.8.3: Paged page table.



A paged virtual memory has the following characteristics:

- The virtual address size is 32 bits.
- The page size is 1024 words.
- Each word is 1 integer.
- Each page table entry needs 4 integers.

- (a) Determine the size of the page table, PT.
- (b) The PT is too large and must be divided into a hierarchy of pages of size 1024 integers. Determine the number of levels in the hierarchy such that no PT exceeds the page size.
- (c) Determine the size of the highest level PT in the hierarchy.

**EXERCISE**

7.8.4: Virtual address translation.



Virtual memory uses paging and segmentation. The segment table is also paged. Non-resident pages are kept on a disk.

Physical memory

0	4	16	52	32	28	48	16
1	8	17	-1	33	48	49	---
2	-2	18	---	34	---	50	---
3	---	19	---	35	60	51	36
4	17	20	---	36	56	52	120
5	16	21	---	37	---	53	16
6	100	22	---	38	---	54	16
7	3	23	---	39	---	55	27
8	-1	24	3	40	24	56	52
9	3	25	-2	41	---	57	47
10	16	26	36	42	---	58	-1
11	0	27	47	43	---	59	3
12	---	28	0	44	40	60	44
13	---	29	12	45	20	61	---
14	---	30	---	46	---	62	---
15	---	31	---	47	---	63	---

Page table
of segment
table

Disk

Block #:	0	1	2	3
Contents:	7 11 14 16	3 16 13 25	-11 4 -2 3	4 106 10 0 ...

- A virtual address has the form (s1, s2, p, w) where each component is 2 bits.
- The page table of the paged segment table starts at address 32.

- A dash denotes a nonexistent page or page table.
 - A positive number, n , in a table means that the page currently resident in memory starting at address n .
 - A negative number, $-n$, in a table means that the page is not resident in memory and must be loaded from the disk. The value of n then gives the disk block containing the missing page. Ex: A -2 in a page table entry indicates that the page is not in main memory and can be found in disk block 2.
- Starting with the page table of the segment table at address 32, trace down the entire memory hierarchy and label each frame as: "page p of the segment table," "page table of segment s ," or "page p of segment s ," where p and s are the actual page and segment numbers. Page frames that remain unmarked are free.
 - Translate virtual address 113 into a physical address and determine the content.
 - Translate the virtual address 68 into a physical address and determine the content. Show all changes to the memory hierarchy as a result of the translation.
 - Show all changes to the memory hierarchy as a result of deleting segment 7.



EXERCISE

7.8.5: Using a TLB.



A TLB is used with the memory system of the previous exercise:

$(s1, s2, p)$	Frame address
0	4
16	52

- Which of the following addresses will find a match in the TLB: 0, 1, 2, 3, 4, 9, 16, 17, 52, 64, 65, 66.
- Which of these addresses cause a page fault? Which result in an error?

7.9 Approximations of the LRU algorithm

The LRU requires the queue of pages to be modified at each memory reference, which makes the algorithm too inefficient in practice. Several algorithms exist that approximate LRU and can be implemented efficiently.

The aging page replacement algorithm

A **referenced bit (r-bit)** is a bit associated with a page and is set automatically by the hardware whenever the page is referenced by any instruction.

An **aging register** is associated with a page and is shifted periodically to the right by 1 bit. Unless the most significant bit is set to 1, the page is aging in the sense that the associated register value is steadily decreasing.

The **aging page replacement algorithm** does not maintain pages sorted in the exact LRU order, but groups together pages referenced during a period of d consecutive references. Each period is represented by 1 bit in a periodically shifting aging register.

Example 7.9.1: Using an aging register.

With d = 50, a 4-bit aging register containing 0100 indicates that the page has not been accessed during the last 50 most recent references but has been accessed at least once during the 50 previous references. A register value 0000 indicates the page has not been accessed during the last 200 references.

The algorithm is invoked periodically after every d reference, where d is a parameter determined experimentally:

- 1. Shift all aging registers to the right by 1, discarding the least significant bit.
- 2. Copy the r-bits into the most significant bits of the aging registers.
- 3. Reset all r-bits to 0.

When a page fault occurs:

- Select the page with the smallest aging register value for replacement.
- If multiple pages have the same smallest value then select one at random.

PARTICIPATION
ACTIVITY

7.9.1: The aging page replacement algorithm.

Page 0

0

0000

0

1000

0

0100

0

1010

Page 1

0

0000

0

0000

0

1000

0

0100

Resident

r

aging

r

aging

r

aging

r

aging

0

0

0

0

0

0

0

0

0

1

0

0

0

1

0

1

0

0

0

0

0

1

0

0

0

1

0

1


0

0

1

0

←



Animation content:

Static figure: A three-column table with headers Resident, r, and aging. Four rows representing Pages 0 to 3 are listed under the Resident column. Each page has an associated r bit and a 4-bit aging register. In the initial state, the r bit and all bits of the aging registers are set to 0. The last 2 columns are replicated 3 times to illustrate changes over time.

Step 1: Initially, pages 0 - 3 are resident. The r-bits and the aging registers all contain zeros.

Step 2: After d references, the r-bits show that pages 0, 2, and 3 have been accessed at least once during the past d references.

The r-bits are 1011 thus reflecting accesses to Pages 0, 2, and 3.

Step 3: All aging registers are shifted to the right. The r-bits are shifted into the most significant positions of the aging registers and then reset to 0.

The aging register for Page 0 now starts with a 1 in the most significant bit (leftmost), followed by 3 zeros. The same applies to Pages 2 and 3. Page 1's register remains at all zeros since its r-bit was 0. All r-bits are now reset to 0.

Step 4: During the next d references, pages 1 and 2 were accessed at least once and so the corresponding r-bits are 1. The r-bits are 0110 thus reflecting accesses to Pages 1 and 2.

Step 5: The aging registers are again shifted to the right. The r-bits are shifted into the most significant register positions and then reset to 0.

The aging registers have the contents: For Pages 0 and 3: 0100. For Page 1, 1000. For Page 2, 1100.

Step 6: During the next d references, pages 0 and 3 are accessed and the same steps of shifting and resetting the r-bits are repeated. The r-bits are 1001 thus reflecting accesses to Pages 0 and 3.

The aging registers have the contents: For Pages 0 and 3: 1010. For Page 1, 0100. For Page 2, 0110.

Step 7: If a page fault occurs then all registers are compared and the page with the smallest value is

selected. Page 1 is replaced with the faulting page. The aging register for Page 1 contains 0100, which is less than both 0110 and 1010.

Animation captions:

1. Initially, pages 0 - 3 are resident. The r-bits and the aging registers all contain zeros.
2. After d references, the r-bits show that pages 0, 2, and 3 have been accessed at least once during the past d references.
3. All aging registers are shifted to the right. The r-bits are shifted into the most significant positions of the aging registers and then reset to 0.
4. During the next d references, pages 1 and 2 were accessed at least once and so the corresponding r-bits are 1.
5. The aging registers are again shifted to the right. The r-bits are shifted into the most significant register positions and then reset to 0.
6. During the next d references, pages 0 and 3 are accessed and the same steps of shifting and resetting the r-bits are repeated.
7. If a page fault occurs then all registers are compared and the page with the smallest value is selected. Page 1 is replaced with the faulting page.

PARTICIPATION ACTIVITY

7.9.2: Contents of an aging register.



The aging registers have 4 bits. The parameter d is 100. A page p is accessed at times 50, 150, and 160.

1) At time 100 the register will contain ____.



- ☐ 1000
- ☐ 1100
- ☐ 0001

2) At time 170 the register of page p will contain ____.



- ☐ 1000
- ☐ 1100
- ☐ 1110

3) At time 200 the register of page p will contain ____.



- ☐ 1000
- ☐ 1100
- ☐ 1110



The aging registers have 4 bits. The parameter d is 100. Page p_1 has not been accessed for more than 400 references and page p_2 has not been accessed for more than 600 references.

When a page fault occurs, the aging algorithm will replace ____.



- ☐ page p_1
- ☐ page p_2
- ☐ either page chosen at random

The second-chance page replacement algorithm

The **second-chance page replacement algorithm** is a coarse-grain approximation of LRU. The algorithm uses the r -bit to divide all pages into only two categories: recently referenced and not recently referenced. A page is selected from the not-recently referenced category.

Similar to FIFO, the algorithm maintains a pointer that cycles through the pages when searching for a page to be replaced. Because of the circular nature, the algorithm is also known as the clock algorithm.

At a page fault, the algorithm repeats the following steps until a page is selected for replacement:

1. If the pointer is at a page p with $r = 0$ then select p and advance the pointer to the next page.
2. Otherwise, reset r to 0, advance the pointer to the next page, and continue the search.



Time t	0	1	2	3	4
RS		2	1	4	0
Frame 0	0 0	0 0	0 0	4 1	▶4 1
Frame 1	1 1	1 1	1 1	1 1	1 0
Frame 2	2 0	2 1	2 1	2 1	2 0
Frame 3	3 0	3 0	3 0	3 0	0 1
Pg fault				*	*

Animation content:

Static figure: A tabular representation of page references over time for an example of the second-chance page replacement algorithm. The table headers include Time *t* with timestamps from 0 to 4, RS (Reference String) with values corresponding to the page references at each timestamp, and Pg fault indicating the occurrence of page faults. There are four rows labeled Frame 0 to Frame 3, each showing two columns: one for page numbers and the other for the reference-bit, *r*, which is either 0 or 1.

Step 1: Pages 0 through 3 are resident in frames 0 through 3. The pointer for starting a new search is at page 0.

Step 2: The bits following the vertical bars are the reference-bits *r*. Only page 1 has been referenced in the past and has the *r*-bit set to 1.

Step 3: At time 1 page 2 is referenced. The *r*-bit is set to 1 but the pointer does not change.

Step 4: At time 2 page 1 is referenced. Since the *r*-bit is already 1, no change is necessary.

Step 5: A page fault occurs at time 3. Since the pointer is at a page with *r* = 0, the page is replaced by page 4 and the pointer is advanced to the next page. RS at time 3 contains 4. An asterisk indicates a page fault. In frame 0, 0|0 is replaced by 4|1 and the pointer advances to frame 1, which contains 1|1.

Step 6: The next page fault starts the search with page 1. Since *r* = 1, page 1 is given a second chance to stay resident. The *r*-bit is reset and the pointer is advanced. RS at time 4 contains 0. An asterisk indicates a page fault. In frame 1, 1|1 is replaced by 1|0 and the pointer advances to frame 2, which contains 2|1.

Step 7: The *r*-bit of the next page 2 is also 1, allowing page 2 to remain resident. In frame 2, 2|1 is replaced by 2|0 and the pointer advances to frame 3, which contains 3|0.

Step 8: The *r*-bit of page 3 is 0 and so page 3 is replaced by the faulting page 0. The pointer cycles back to page 4. In frame 3, 3|0 is replace by 0|1 and the pointer advances to frame 0.

Animation captions:

- 1. Pages 0 through 3 are resident in frames 0 through 3. The pointer for starting a new search is at page 0.

2. The bits following the vertical bars are the reference-bits r . Only page 1 has been referenced in the past and has the r -bit set to 1.
3. At time 1 page 2 is referenced. The r -bit is set to 1 but the pointer does not change.
4. At time 2 page 1 is referenced. Since the r -bit is already 1, no change is necessary.
5. A page fault occurs at time 3. Since the pointer is at a page with $r = 0$, the page is replaced by page 4 and the pointer is advanced to the next page.
6. The next page fault starts the search with page 1. Since $r = 1$, page 1 is given a second chance to stay resident. The r -bit is reset and the pointer is advanced.
7. The r -bit of the next page 2 is also 1, allowing page 2 to remain resident.
8. The r -bit of page 3 is 0 and so page 3 is replaced by the faulting page 0. The pointer cycles back to page 4.

**PARTICIPATION
ACTIVITY**

7.9.5: Replacement choices made by the second-chance algorithm.



Pages 0 - 2 reside in frames 0 - 2. All r -bits are 0 and the pointer is at page 0:

	Page r
Frame 0	0 0
Frame 1	1 0
Frame 2	2 0

Assume the reference string $RS = \{3, 2, 1, 0\}$.

1) The first page fault will replace page

_____.

- ☐ 0
- ☐ 1
- ☐ 2



2) The second page fault will replace page

_____.

- ☐ 0
- ☐ 1
- ☐ 2



3) With LRU, the second page fault would replace page ____.

- ☐ 1
- ☐ 2
- ☐ 3

The third-chance page replacement algorithm

The third-chance algorithm is a refinement of the second-chance algorithm. The algorithm takes advantage of both the r-bit and m-bit associated with every page. Since replacing a modified page is more costly than replacing an unmodified page, the algorithm gives modified pages an additional chance to remain resident.

The **third-chance page replacement algorithm**, also known as the **not-recently-used page replacement algorithm** (NRU), is a coarse-grain approximation of LRU, which divides pages into 4 categories based on the 4 possible combination of the r- bit and the m-bit.

At a page fault, the algorithm repeats the following steps until a page is selected for replacement:

- 1. If the pointer is at a page p with r = 0 and m = 0 then select p and advance the pointer to the next page.
- 2. Otherwise, reset the bits according to the following table, advance the pointer to the next page, and continue the search.

11 → 01	modified page
01 → 00	
10 → 00	unmodified page

A third bit is needed to remember if a page has been modified and thus needs to be written back to disk when selected for replacement.



Time t	0	<u>1</u>	2	3	<u>4</u>	<u>5</u>
RS		0	4	<u>0</u>	<u>2_w</u>	<u>1</u>
Frame 0	0 00	0 10	0 00	0 10	0 10	0 00
Frame 1	1 00	1 00	4 10	4 10	4 10	4 00
Frame 2	2 10	2 10	2 10	2 10	2 11	2 01
Frame 3	3 10	3 10	3 10	3 10	3 10	3 00
Pg fault			*		*	

Animation content:

Static figure: A table illustrating the third-chance page replacement algorithm.

The table headers include Time t with timestamps from 0 to 5, RS (Reference String) with values corresponding to the page references at each timestamp, and Pg fault indicating the occurrence of page faults. There are four rows labeled Frame 0 to Frame 3, each showing two columns: one for page numbers and the other for two bits: the reference-bit, r, and the modified-bit, m, both of which are either 0 or 1.

Step 1: Pages 0 - 3 are resident. Pages 2 and 3 have been referenced in the past and have the r-bit set to 1. No page has been modified. The pointer for starting a new search is at page 0.

The rm-bits at time 0 are: 00 for pages 0 and 1, and 10 for pages 2 and 3.

Step 2: At time 1, page 0 is referenced. Since the page is resident, only the r-bit is set.

The rm-bits of page 0 in frame 0 change from 00 to 10.

Step 3: The page fault at time 2 starts the search at page 0. Since the 2 bits are not 00, page 0 gets a second chance to remain resident. The r-bit is reset and the pointer moves to the next page.

Page 4 is referenced at time 2. An asterisk at time 2 marks the page fault. The rm-bits of page 0 change from 10 to 00. The pointer moves to the next frame 1.

Step 4: Page 1's r/m-bits are 00. Thus page 1 is replaced by page 4 and the pointer is advanced to the next page. The entry 1|00 in frame 1 is replaced by 4|10. The pointer moves to the next frame 2.

Step 5: At time 3, page 0 is referenced again, which sets the r-bit to 1. The r/m-bits of all pages are now 10. The rm-bits 00 in frame 0 are replaced by 10.

Step 6: At time 4, page 2 is modified. The w-subscript indicates a write operation. The m-bit is set. Page 2 is referenced at time 4. A w-mark indicates modification. The rm-bits of page 2 change from

10 to 11.

Step 7: The second page fault at time 5 starts the search with page 2. The r/m-bits are both set, giving page 2 an additional chance to remain resident. The r-bit is reset first.

Page 1 is referenced at time 5. An asterisk at time 5 marks the page fault. The rm-bits of page 2 change from 11 to 01. The pointer moves to the next frame 3.

Step 8: Next, pages 3, 0, and 4 are examined in sequence. Each time the r-bit is reset and the pointer advanced to the next page. The following sequence of changes occurs:

In Frame 3: 3|10 -> 3|00. Pointer advances to frame 0.

In Frame 0: 0|10 -> 0|00. Pointer advances to frame 1.

In Frame 1: 4|10 -> 4|00. Pointer advances to frame 2.

Step 9: Since no page with bits 00 has been found after a complete cycle, the search continues with page 2. The m-bit is reset and the pointer advanced to page 3, which is replaced by page 1.

In Frame 2: 2|01 -> 2|00. Pointer advances to frame 3.

In Frame 3: page 3 is replaced by page 1 and the r-bit is set: 3|00 -> 1|10.

Animation captions:

1. Pages 0 - 3 are resident. Pages 2 and 3 have been referenced in the past and have the r-bit set to 1. No page has been modified. The pointer for starting a new search is at page 0.
2. At time 1, page 0 is referenced. Since the page is resident, only the r-bit is set.
3. The page fault at time 2 starts the search at page 0. Since the 2 bits are not 00, page 0 gets a second chance to remain resident. The r-bit is reset and the pointer moves to the next page.
4. Page 1's r/m-bits are 00. Thus page 1 is replaced by page 4 and the pointer is advanced to the next page.
5. At time 3, page 0 is referenced again, which sets the r-bit to 1. The r/m-bits of all pages are now 10.
6. At time 4, page 2 is modified. The w-subscript indicates a write operation. The m-bit is set.
7. The second page fault at time 5 starts the search with page 2. The r/m-bits are both set, giving page 2 an additional chance to remain resident. The r-bit is reset first.
8. Next, pages 3, 0, and 4 are examined in sequence. Each time the r-bit is reset and the pointer advanced to the next page.
9. Since no page with bits 00 has been found after a complete cycle, the search continues with page 2. The m-bit is reset and the pointer advanced to page 3, which is replaced by page 1.



1) A page with the r/m-bits set to 01 has been ____.



- ☐ modified but not referenced
- ☐ referenced but not modified
- ☐ referenced and modified

2) When a page with the bits 00 is modified, the bits are set to ____.



- ☐ 10
- ☐ 01
- ☐ 11

3) With 3 frames, $RS = \{0\ 1\ 2\ 1\ 3\}$ where the first reference modifies the page 0, and the pointer at page 0, the first page fault will replace page ____.



- ☐ 0
- ☐ 1
- ☐ 2



726670.4301284.qx3zqy7

Start

A system uses the second-chance page replacement algorithm.

At time i , the four page frames contain the page numbers and r-bits shown below. The pointer is at the page to be considered next, page 6.

Time t	...	i	...
Frame 0		4 0	
Frame 1		7 1	
Frame 2		→6 0	

At time $i + 1$, page p is referenced. Determine which of the values change (if any) when each page is referenced.

	Page number	r-bit	Pointer
$p = 8$	☐	☐	☐
$p = 4$	☐	☐	☐
$p = 5$	☐	☐	☐

1	2	3	4	5	6
---	---	---	---	---	---

Check

Next



EXERCISE

7.9.1: Using aging registers.



A memory consists of 4 page frames, currently holding pages 0 - 3. A 6-bit aging register is associated with each frame. The table shows which pages were referenced during each of 8 consecutive periods d .

Period d	0	1	2	3	4	5	6	7
Pages referenced	3	1 2	0	0 3	0	2	2	1

(a) Determine which page will be replaced if a page fault occurs at the end of last d period.



EXERCISE

7.9.2: Comparing page replacement algorithms.



Physical memory consists of 4 page frames, initially all empty. The following reference string is processed:

0 1 4 0 2 3 0 1 0 2 3 4 2 3

- (a) Show which pages are resident under the second chance page replacement algorithm. Indicate when page faults occur.
- (b) Assume that references to page 1 are write references (modifying page 1) and all others are read references. Show which pages are resident under the third chance page replacement algorithm. Indicate when page faults occur.